



Różne sposoby tworzenia modeli Machine Learning w Azure – od graficznych po zarządzanie w modelami w języku Python. W tle klastry obliczeniowe (z GPU).

Tomasz Kopacz  
CEE ISV Technical Recruit Lead  
Microsoft



## Security & Management

 Security Center

 Portal

 Azure Active Directory

 Azure AD B2C

 Multi-Factor Authentication

 Automation

 Scheduler

 Key Vault

 Store/ Marketplace

 VM Image Gallery & VM Depot

## Media & CDN



## Integration



## Application Platform



## Compute Services



## Developer Services



## Data



SQL, MySQL,  
PostgreSQL  
Database



SQL Data  
Warehouse



CosmosDB



SQL Server  
Stretch Database



Redis Cache



Storage  
Tables



Azure  
Search

## Intelligence



Cognitive Services



Bot Framework



## Compute



## Storage



## Datacenter Infrastructure (36 + 6 Regions)

## Hybrid Cloud

 Azure AD Health Monitoring

 AD Privileged Identity Management

 Domain Services

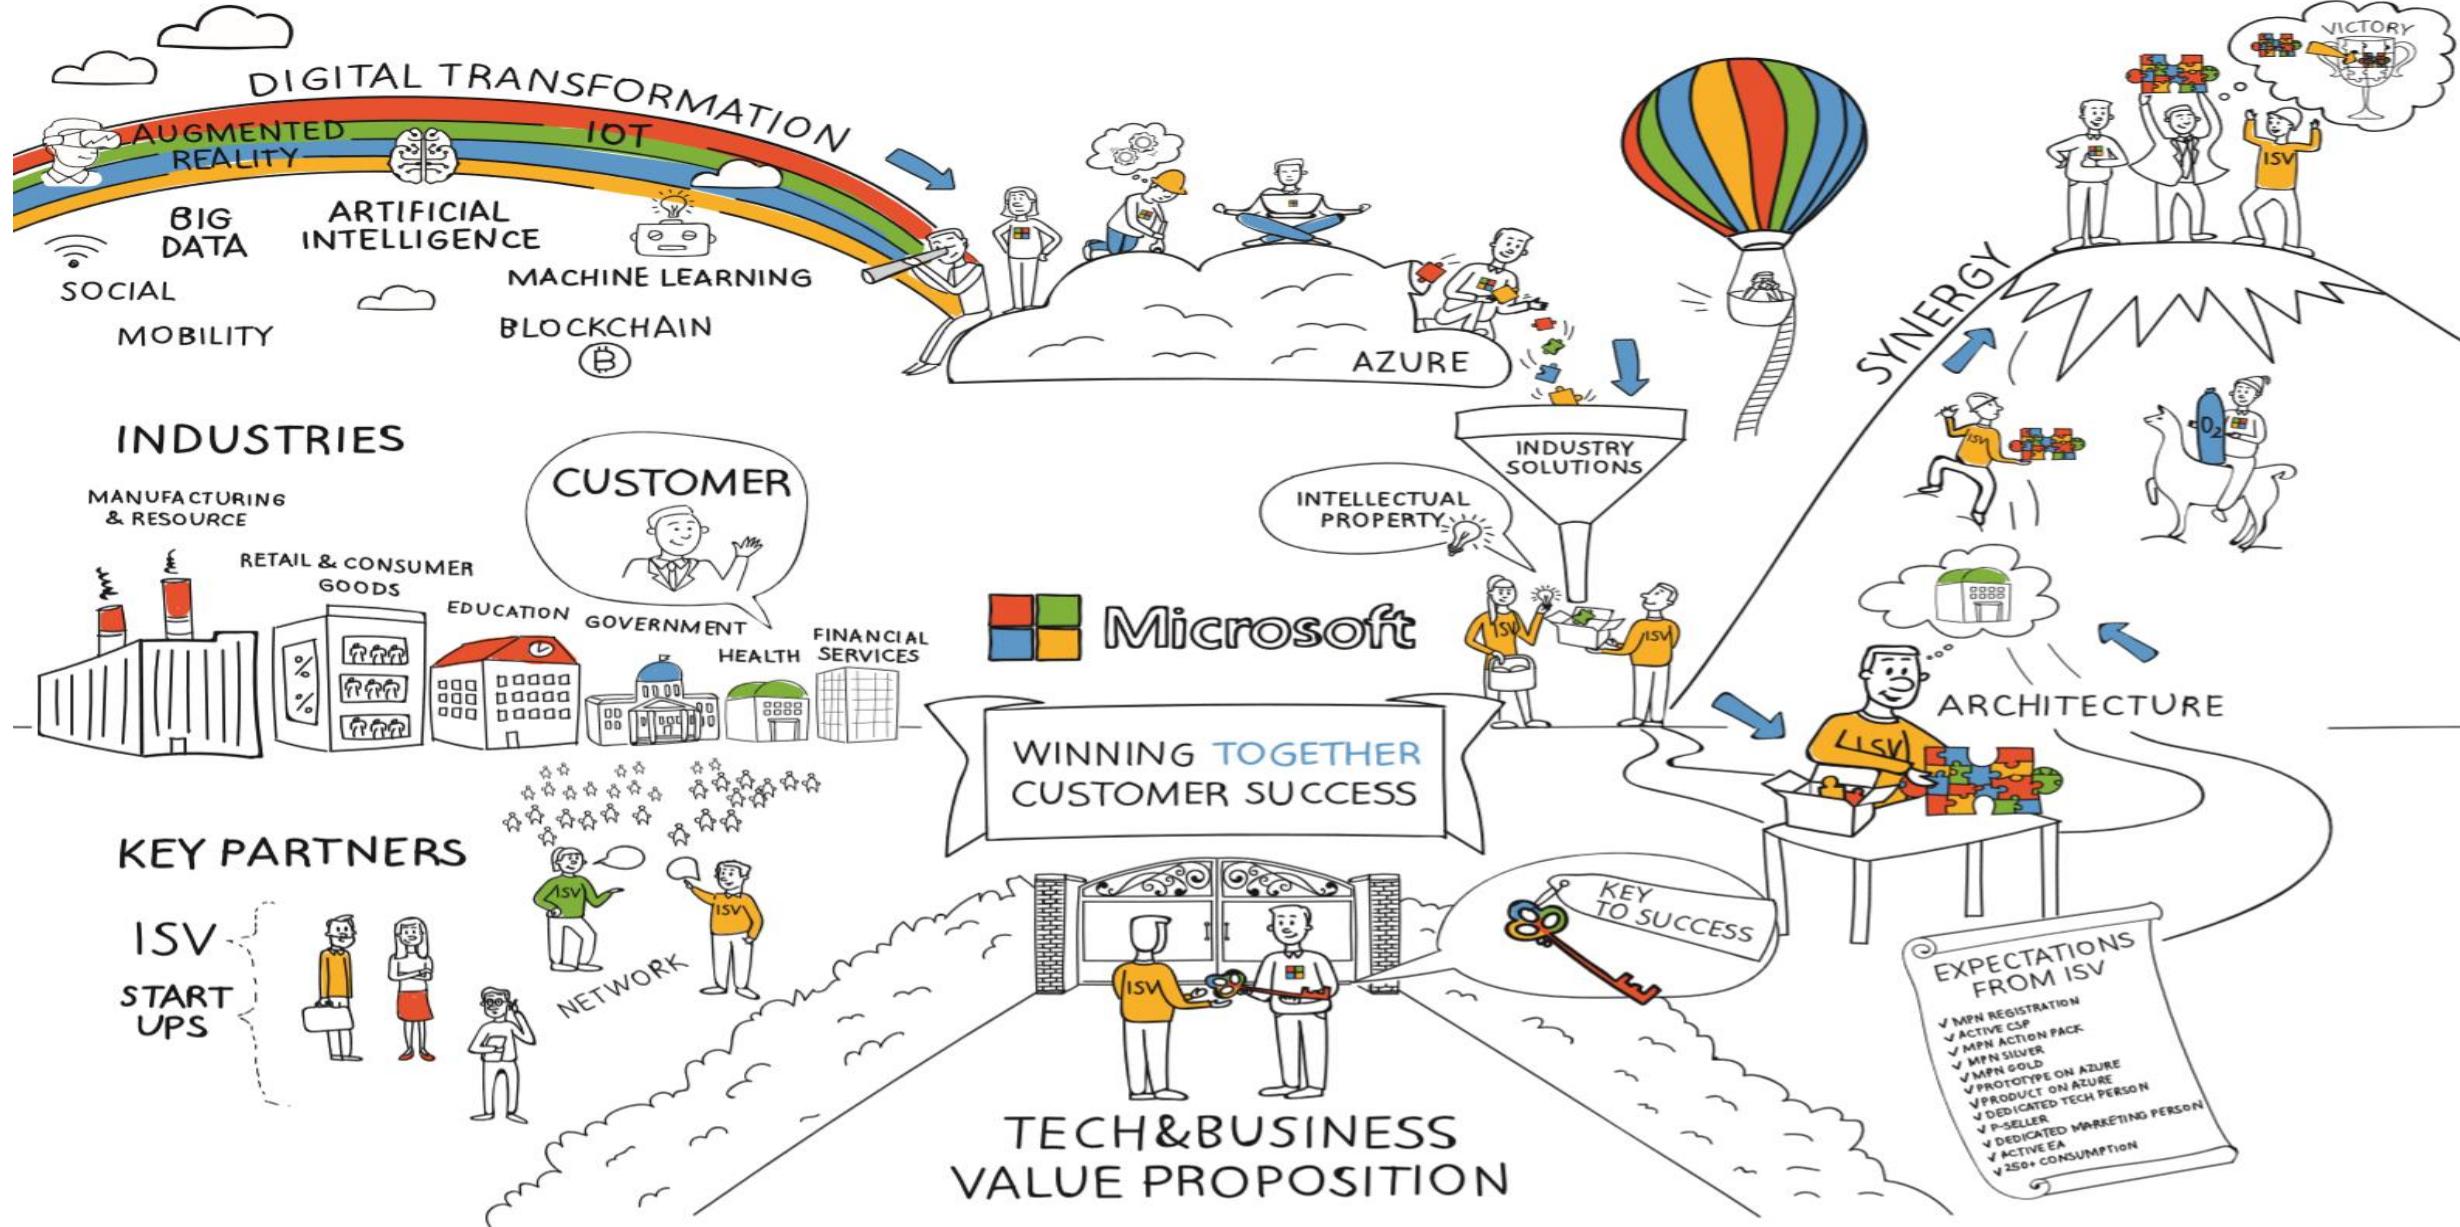
 Backup

 Operational Analytics

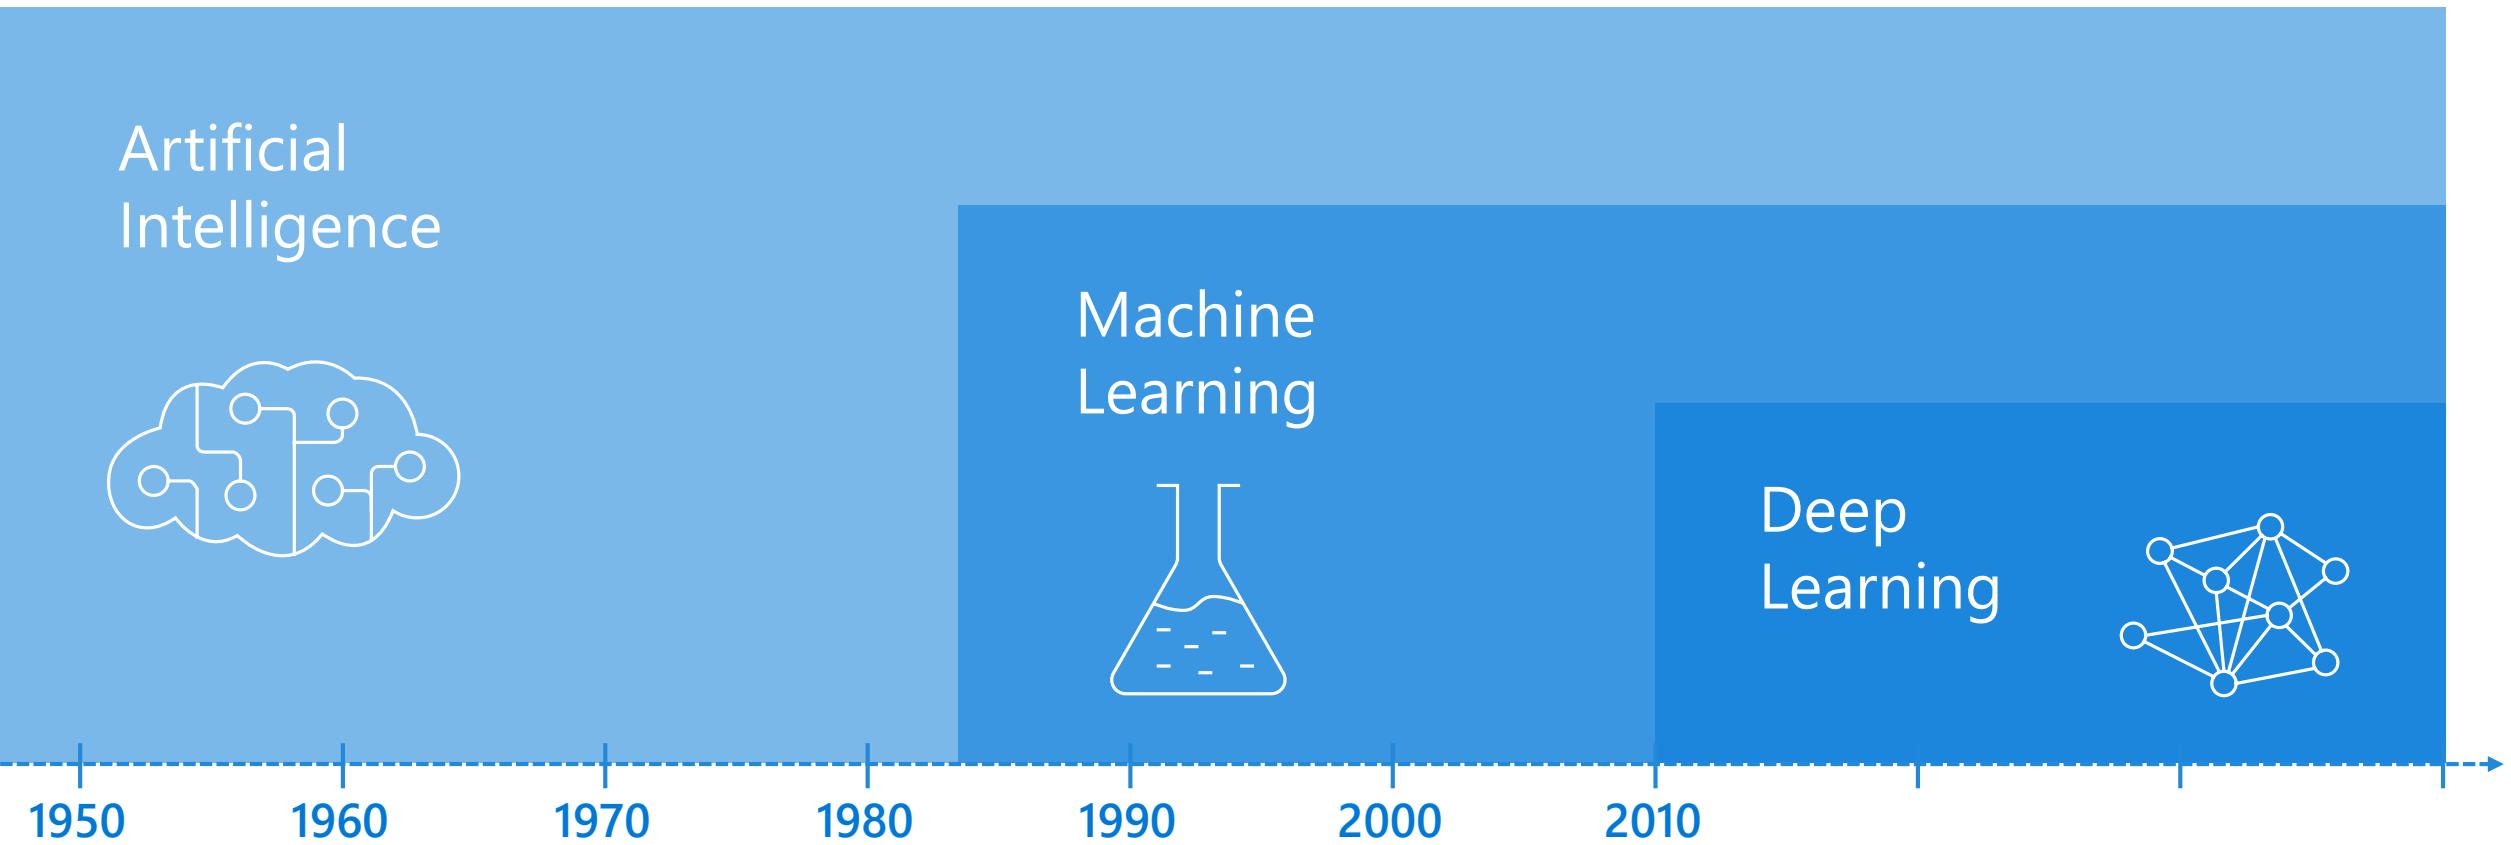
 Import/Export

 Azure Site Recovery

 StorSimple



# AI, Machine Learning and Deep Learning



# Why Learn?

Learn it when you can't code it

(e.g. Recognizing Speech/image/gestures)

Learn it when you can't scale it

(e.g. Recommendations, Spam & Fraud detection)

Learn it when you have to adapt/personalize

(e.g. Predictive typing)

Learn it when you can't track it

(e.g. AI gaming, robot control)

**Learn it when build formal, classic algorithm is too complicated**

Loaded settings from C:\Users\akapo0r\Documents\AirSim\settings.json

Default config: PhysXCar4x4

Press F1 to see help

API server started at (default)

Reverse: (API)

Throttle: (API)

Steering: (API)

Footbrake: (API)

Speed: 1 km/h

Gear: 1



# Future of Software

Program Logic

+

Artificial Intelligence

# **SW Development**

Data Schemas

Functions & Code

Debugger to process

Use telemetry to inform

# **AI Development**

Data Curation + Data Labeling

Machine Learning Models

A/B Testing, Experimentation

Telemetry driven learning

# Mathematics used in AI (PART OF)

## Linear Algebra

Vectors: definition, scalars, addition, scalar multiplication, inner product(dot product)...

Matrices: definition, addition, transpose, scalar multiplication, matrix multiplication...

Eigenvalues & eigenvectors: concept, intuition, significance, how to find

Principle component analysis, Singular value decomposition

## Calculus

Functions, Scalar derivative: intuition, common rules of differentiation, chain rule, partial derivatives

Gradient: concept, intuition, properties, directional derivative, Vector and matrix calculus: how to find derivative of ...

Gradient algorithms: local/global maxima and minima, saddle point, convex functions, ...

## Probability

Basic rules and axioms: events, sample space, frequentist approach, dependent and independent events, conditional ...

Random variables: continuous and discrete, expectation, variance, distributions- joint and conditional

Bayes' Theorem, MAP, MLE, Popular distributions: binomial, bernoulli, poisson, exponential, gaussian, Conjugate priors

## Miscellaneous

Information theory: entropy, cross-entropy, KL divergence, mutual information

Markov Chain: definition, transition matrix, stationarity

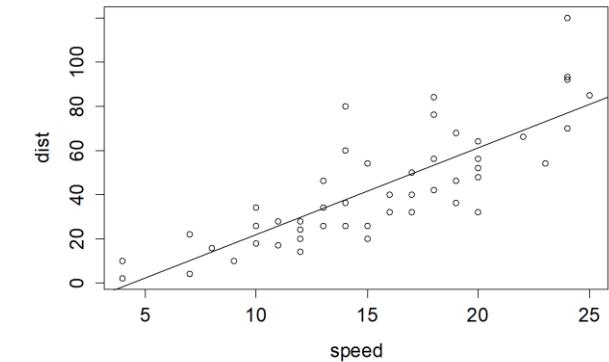
# AI / Machine learning techniques

## Supervised learning (task driven)

Training data includes examples of correct solutions to the problem

Find/generalize relationship between the input data and solution

Example: regression, classification

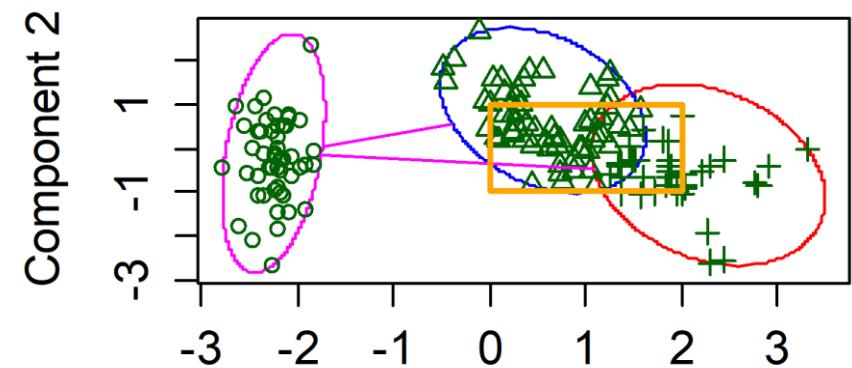


## Unsupervised learning (data driven)

Points are not associated with known output values

Creates a model that learns inherent structure in training data

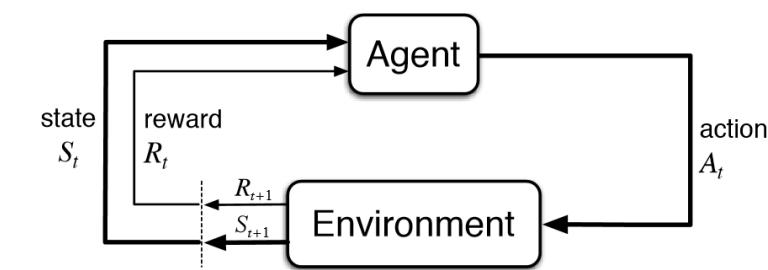
Example: clustering, time-series



## Reinforcement learning (react)

Observations gathered from the interaction with the environment to take actions that would maximize the reward or minimize the risk

Example: Q-Learning, SARSA, DQN



# What are the “neural networks”

This is just a special kind of algorithm!

We can do clustering using DNN or regression or time series ...

“Classification”

Artificial Neural Network

Deep Neural Network (generic type neural network with many hidden layers)

Deep Belief networks

Recurrent neural networks (robot control, music composition, grammar learning, time series)

Convolutional neural network (mainly cognitive vision services; [object detection](#))

...

See: [https://en.wikipedia.org/wiki/Types\\_of\\_artificial\\_neural\\_networks](https://en.wikipedia.org/wiki/Types_of_artificial_neural_networks)

# DNN - „Background“

The activating function sums up the "input signals" and reports the output signal

The initial weights are random

After a given learning phase, the weights are corrected to reduce the error

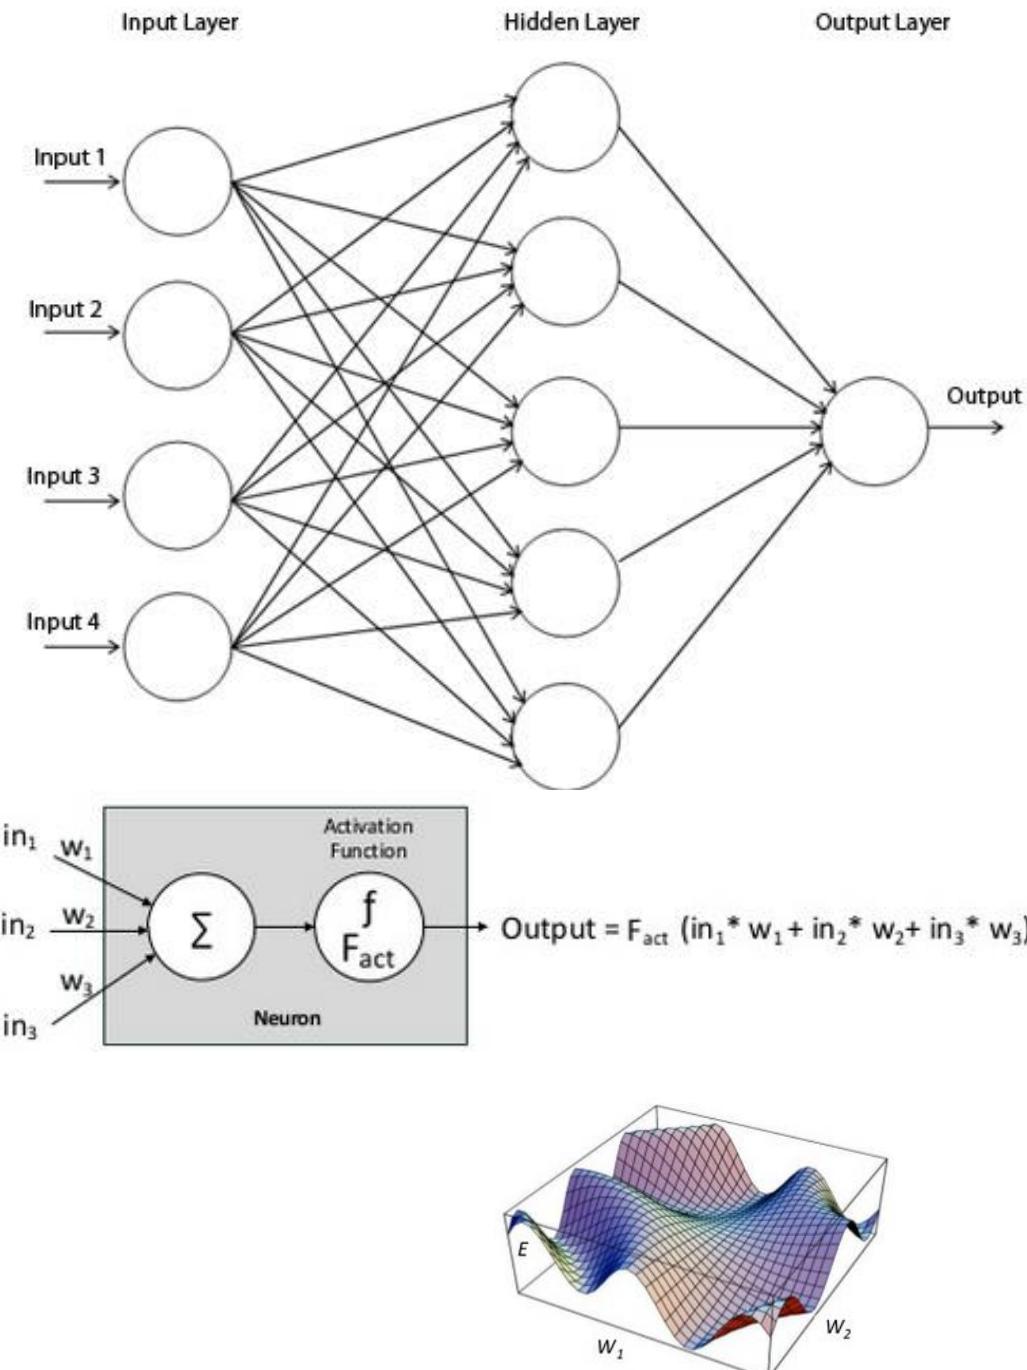
Changes in weights are controlled by the current learning rate

CNTK: Stochastic Gradient Descent, optimization alghoritm, change LR

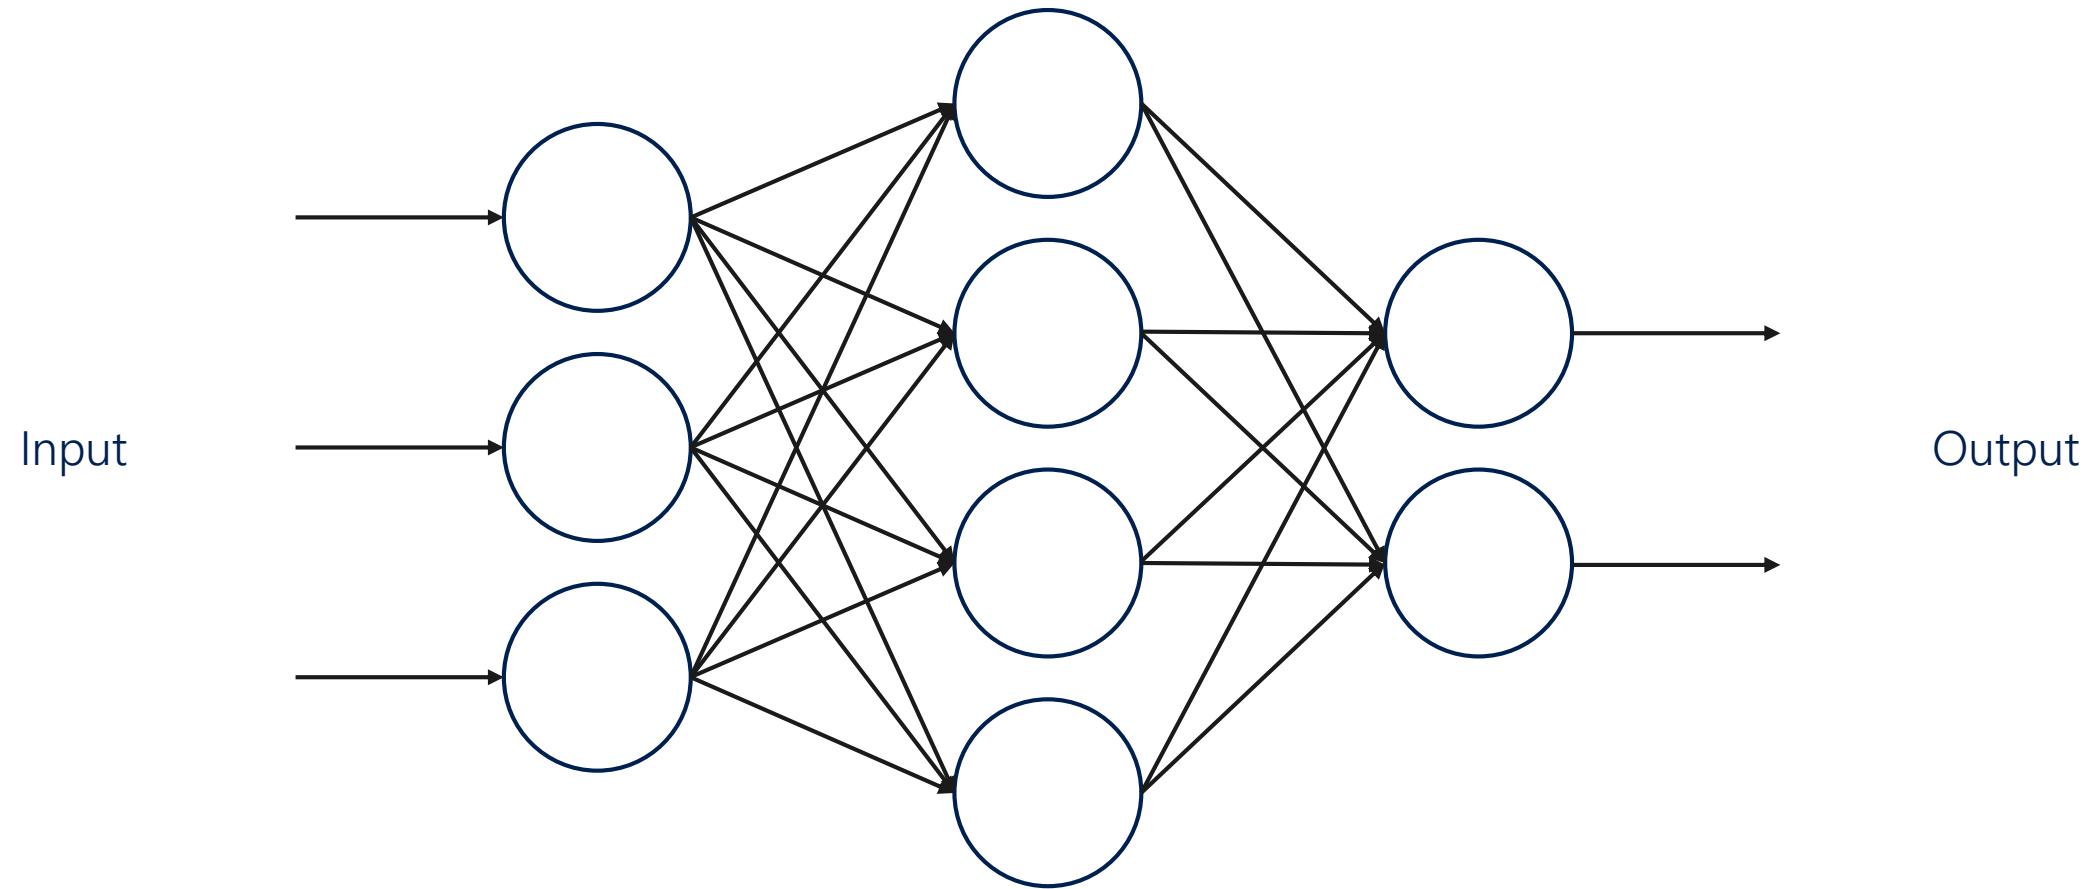
([1 bit](#) to scale out)

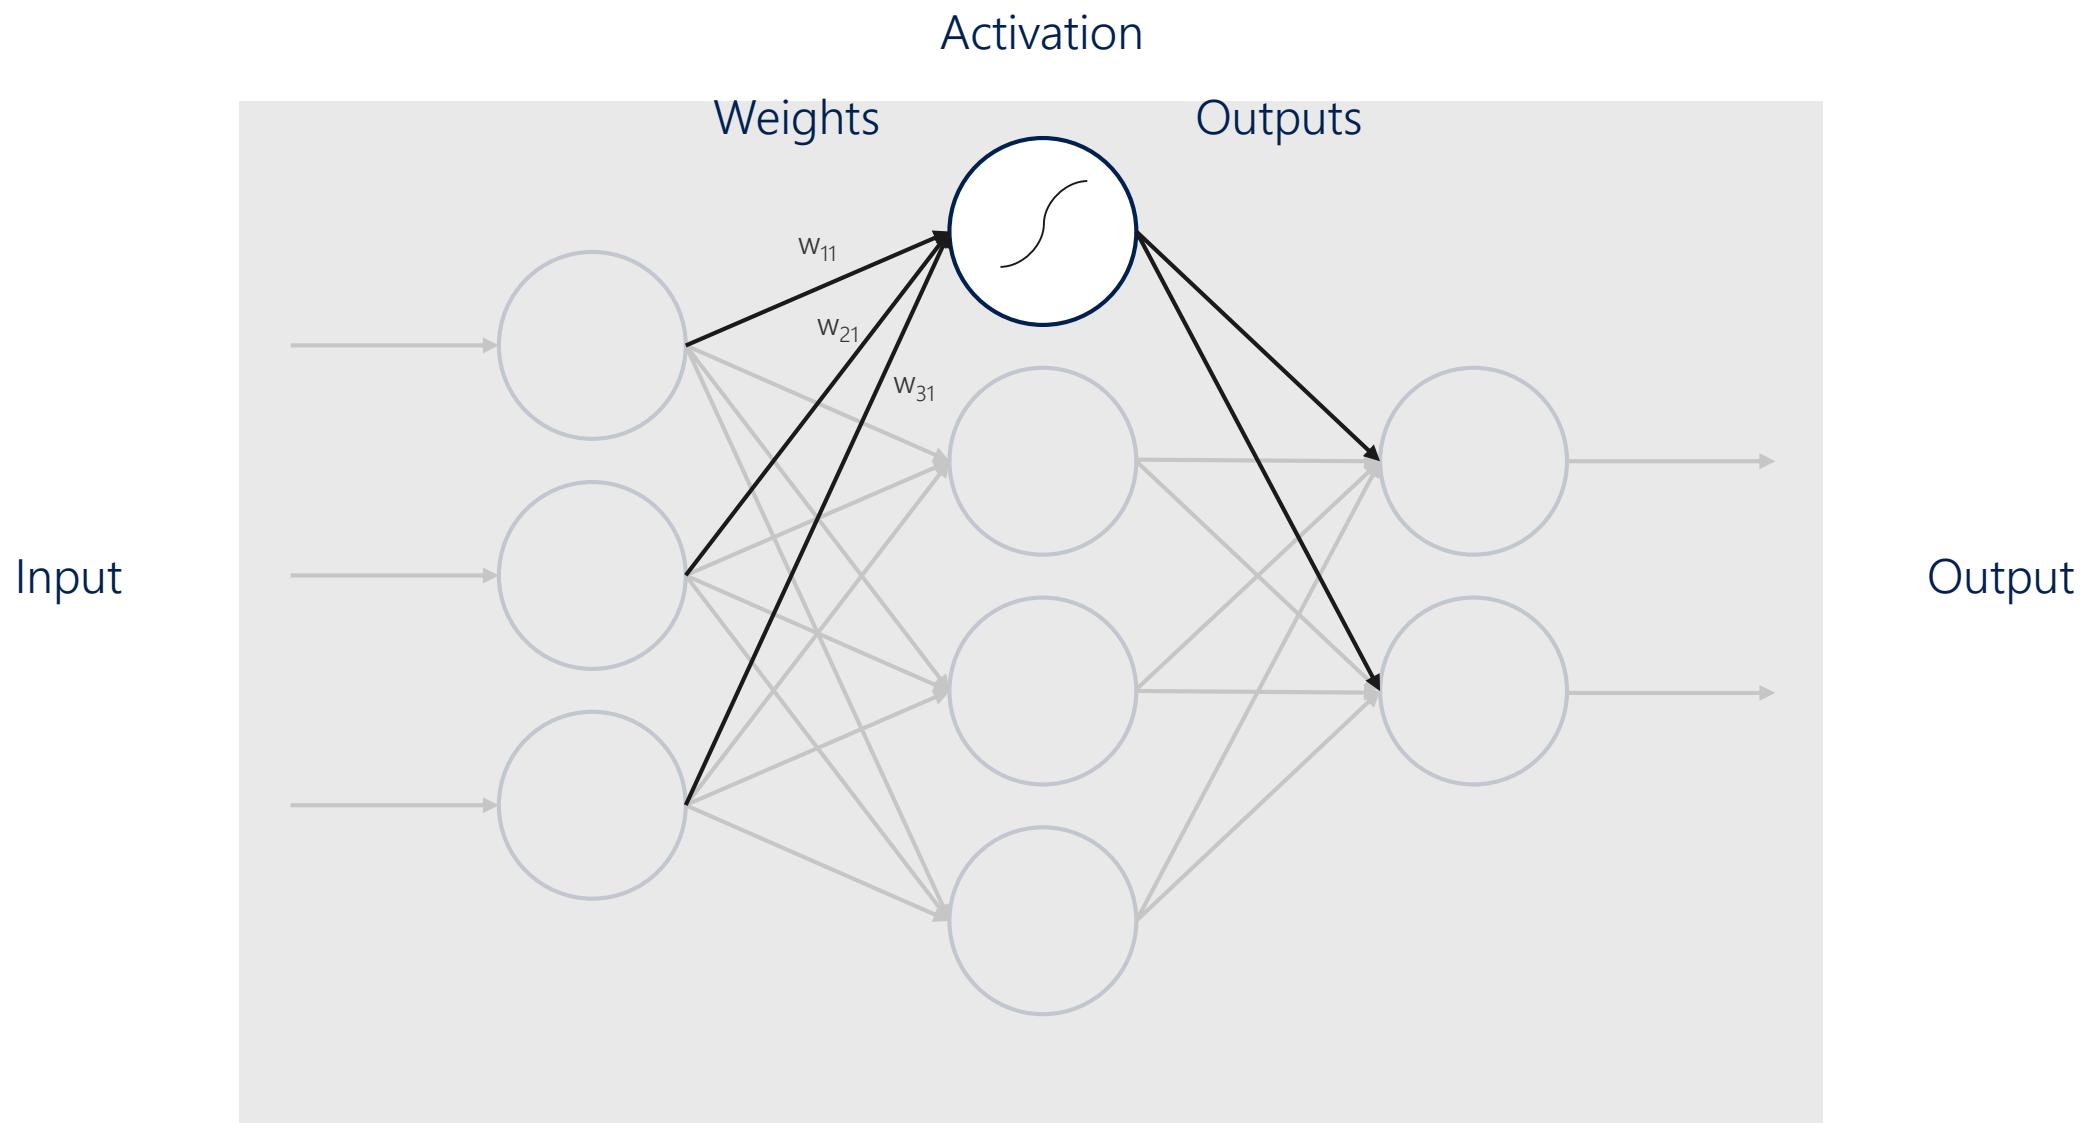
There are many local maximums!

For many weight attributes they create a multidimensional space

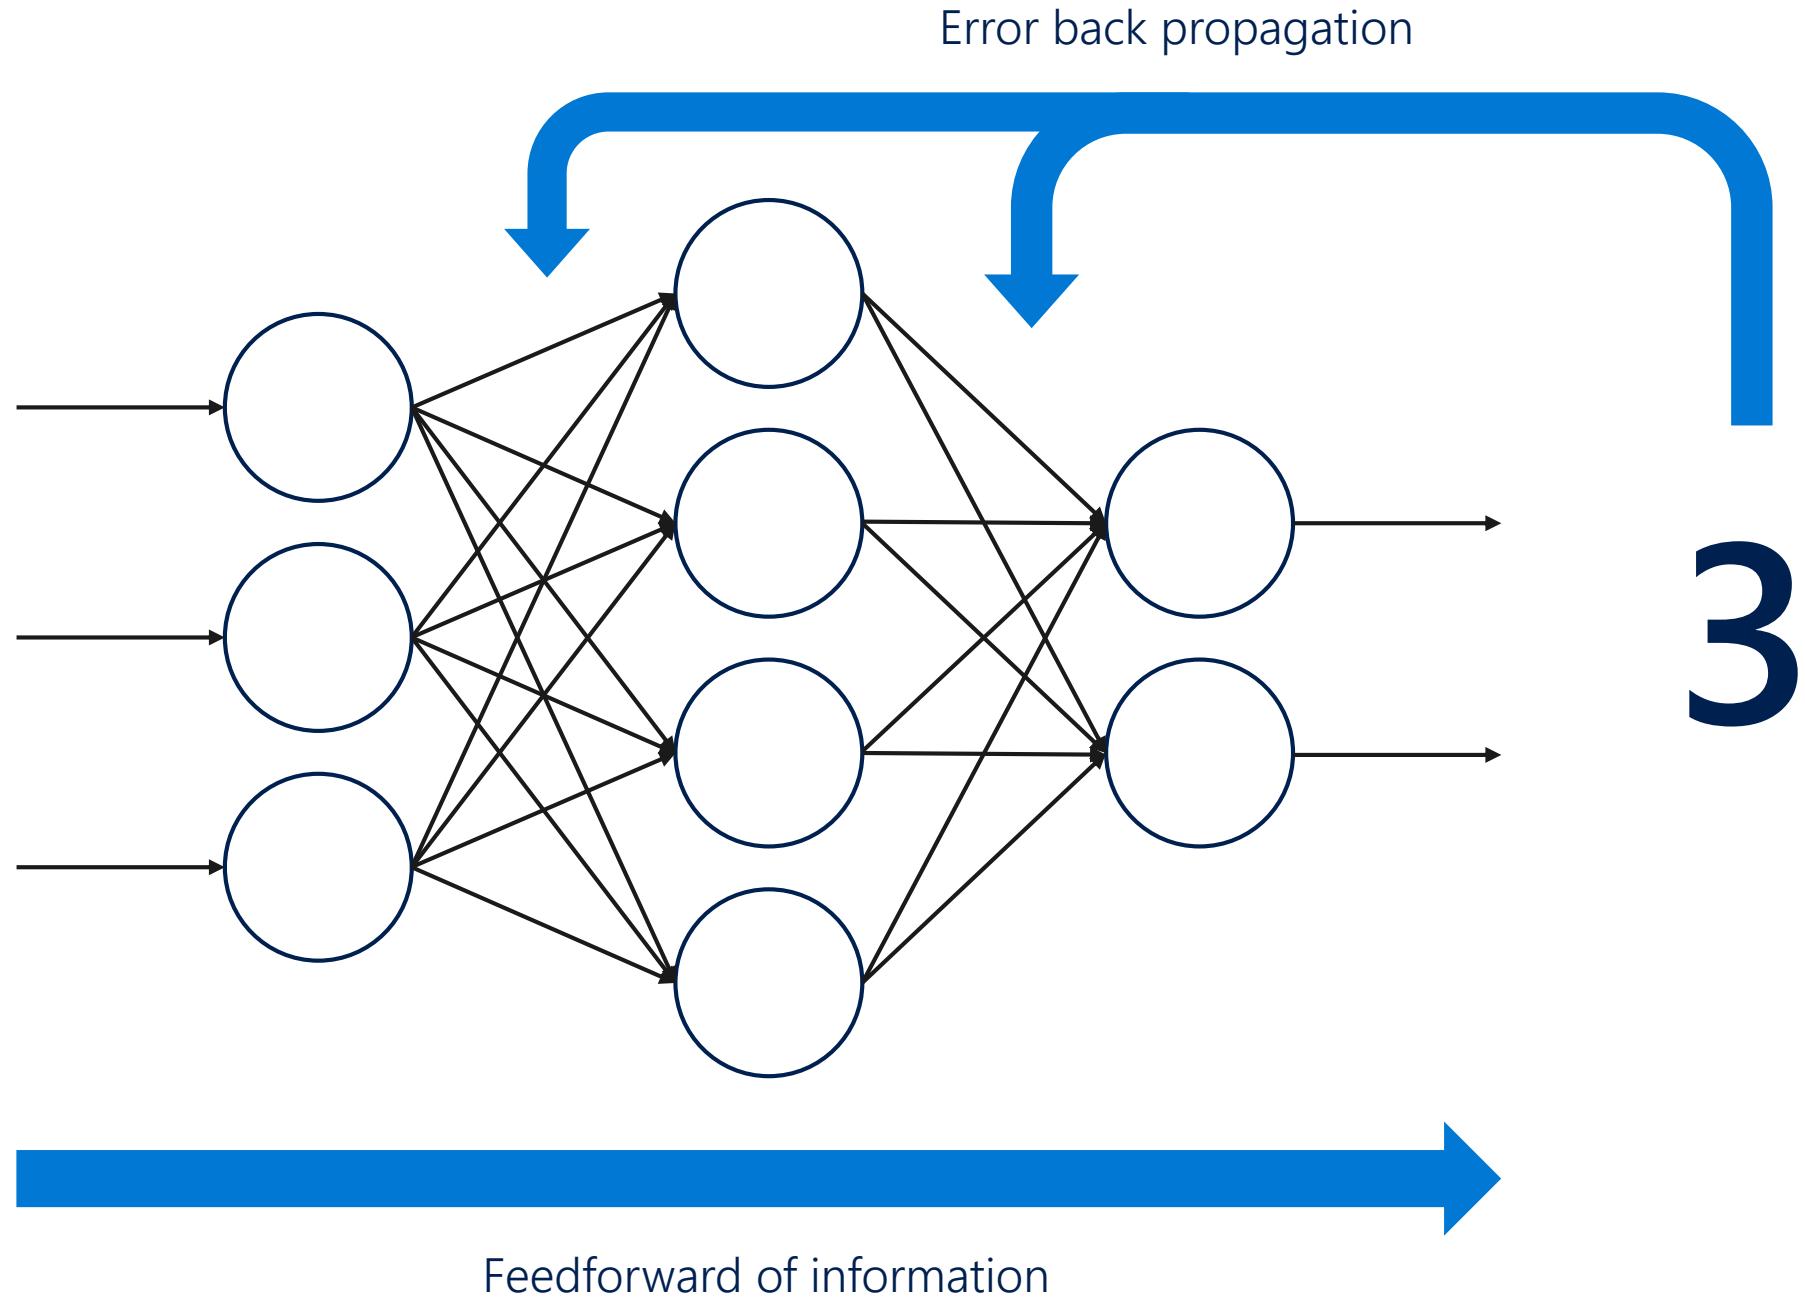


Source: Oberhofer, W., T. Zimmerer, and D.-K. T. Zimmerer (1996). Wie Künstliche Neuronale Netze lernen, p. 16





2



Feedforward of information

3

Error back propagation

# Can Machine Learning Help Me?

- Automated prediction is core
- Lots of past data already available
- Magic numbers in current prediction system

Yes

- Prediction is small part of experience
- No past data available
- Many business-rules govern the experience
- Predictions do not have a predictable pattern

No

# Choose algorithm?



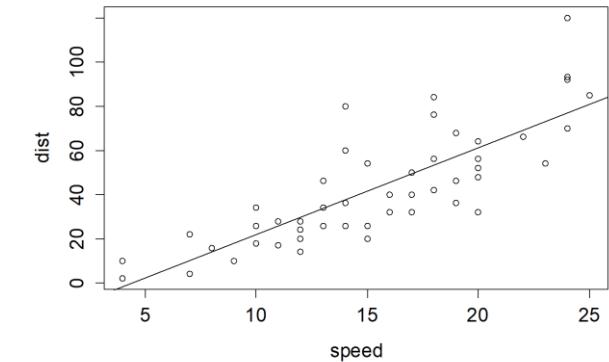
# Machine learning techniques

## Supervised learning (task driven)

Training data includes examples of correct solutions to the problem

Find/generalize relationship between the input data and solution

Example: regression, classification

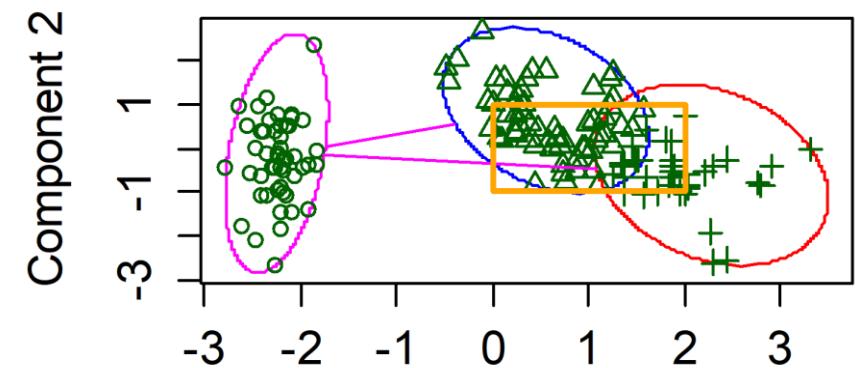


## Unsupervised learning (data driven)

Points are not associated with known output values

Creates a model that learns inherent structure in training data

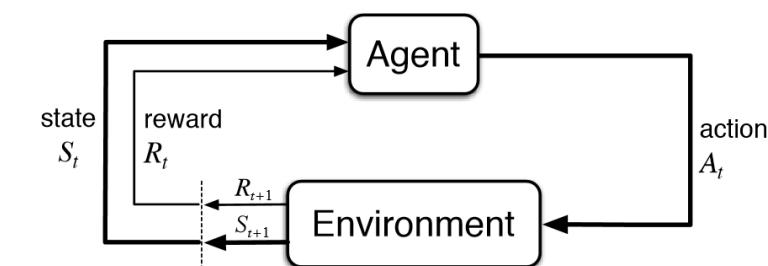
Example: clustering, time-series



## Reinforcement learning (react)

Observations gathered from the interaction with the environment to take actions that would maximize the reward or minimize the risk

Example: Q-Learning, SARSA, DQN



# Develop a Model - Classification

## Classification

Scenarios:

Which customer are more likely to buy, stay, leave (churn analysis)

Which transactions/actions are fraudulent

Which quotes are more likely to become orders

Recognition of patterns: speech, speaker, image, movement, etc.

Algorithms: Boosted Decision Tree, Decision Forest, Decision Jungle, Logistic Regression, SVM, ANN

# Develop a Model - Regression

## Regression

Scenarios:

Stock prices prediction

Sales forecasts

Premiums on insurance based on different factors

Quality control: number of complaints over time based on product specs, utilization, etc.

Workforce prediction

Workload prediction

Algorithms: Bayesian Linear, Linear Regression, Ordinal Regression, ANN

Boosted Decision Tree, Decision Forest

# Develop a Model - Clustering

## Clustering

### Scenarios:

Customer segmentation: divide a customer base into groups of individuals that are similar in specific ways relevant to marketing, such as age, gender, interests, spending habits, etc.

Market segmentation

Quantization of all sorts, such as, data compression, color reduction, etc.

Pattern recognition

### Algorithms: K-means

Sometimes: First, Principal Component Analysis to reduce number of dimensions

# Develop a Model - Recommendations

## Recommendations

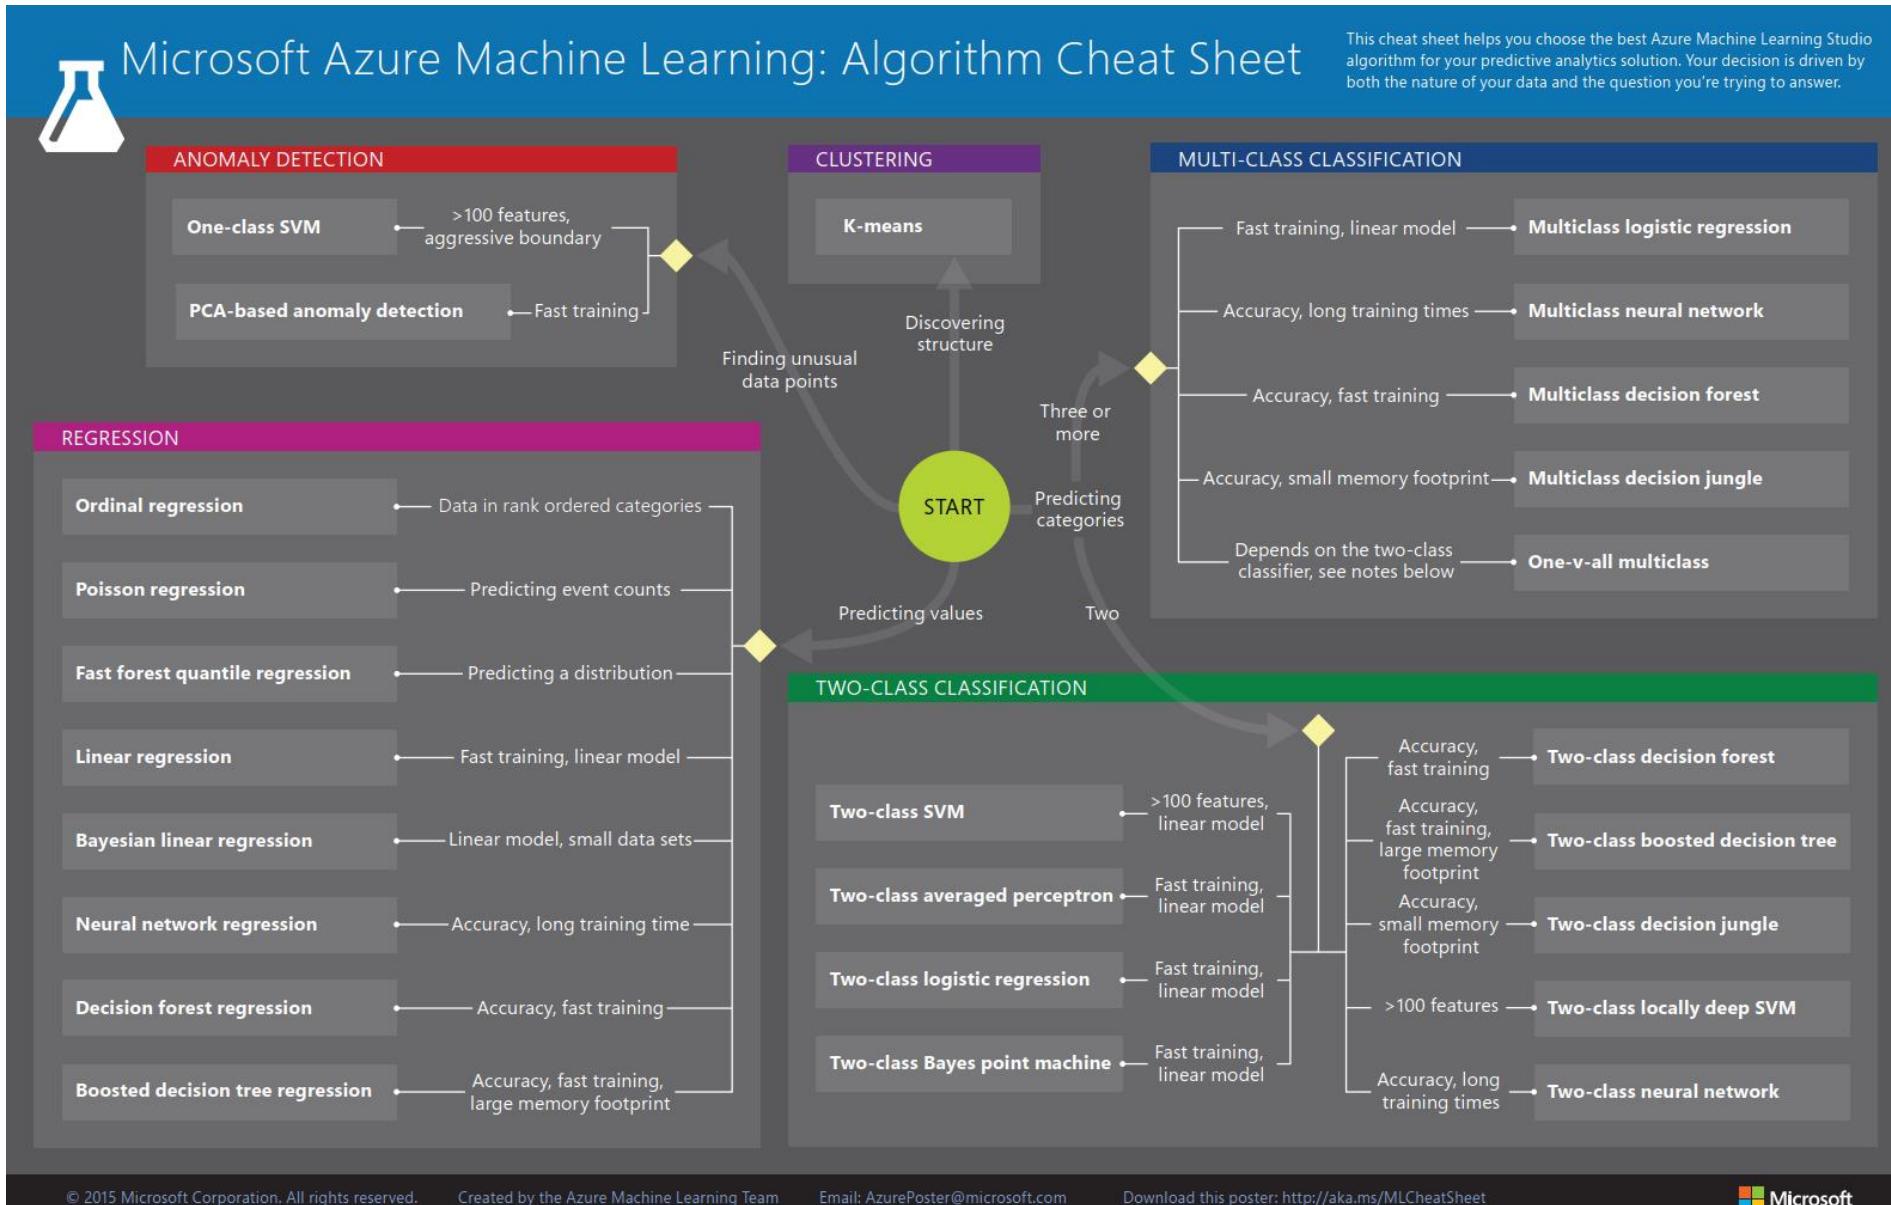
Scenarios:

Product recommendations

Netflix video recommendations

Algorithms: Matchbox

<http://aka.ms/MLCheatSheet>



# Validate model

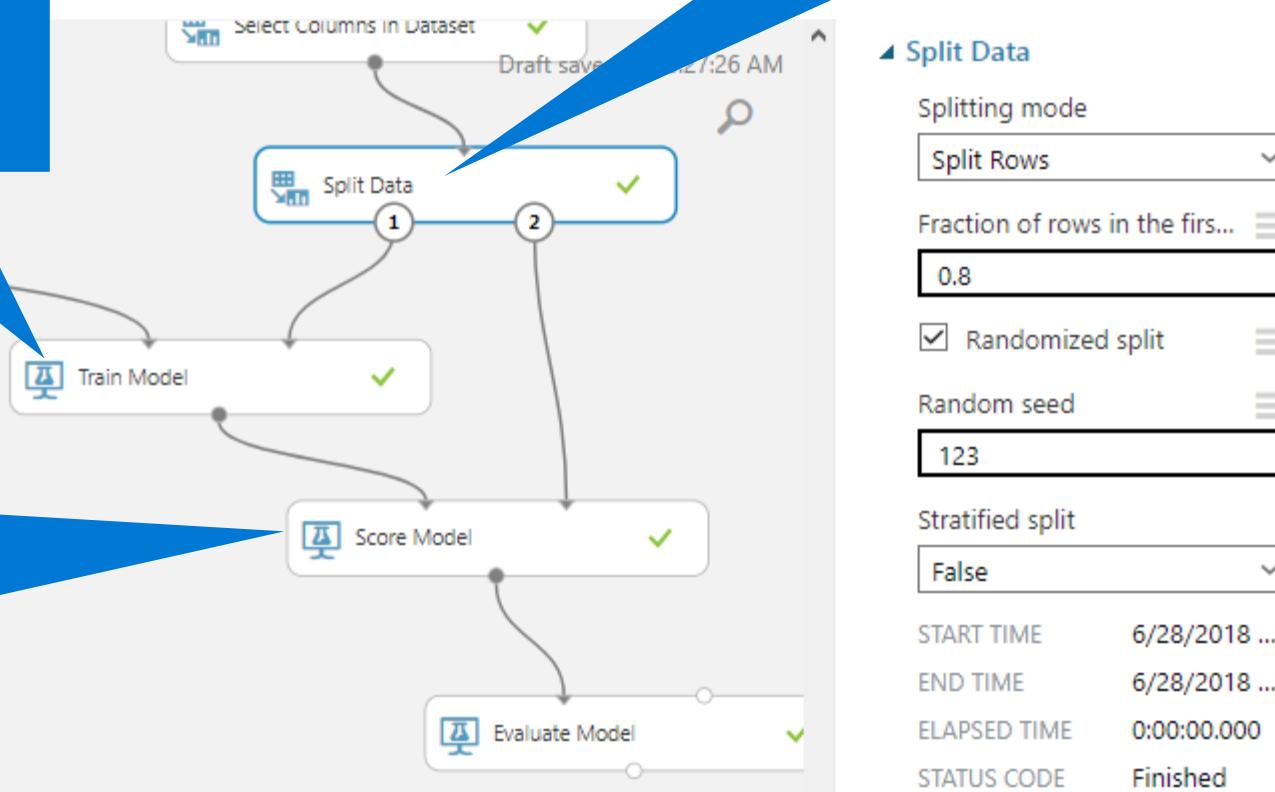


# Testing

Train model (using attached algorithm)

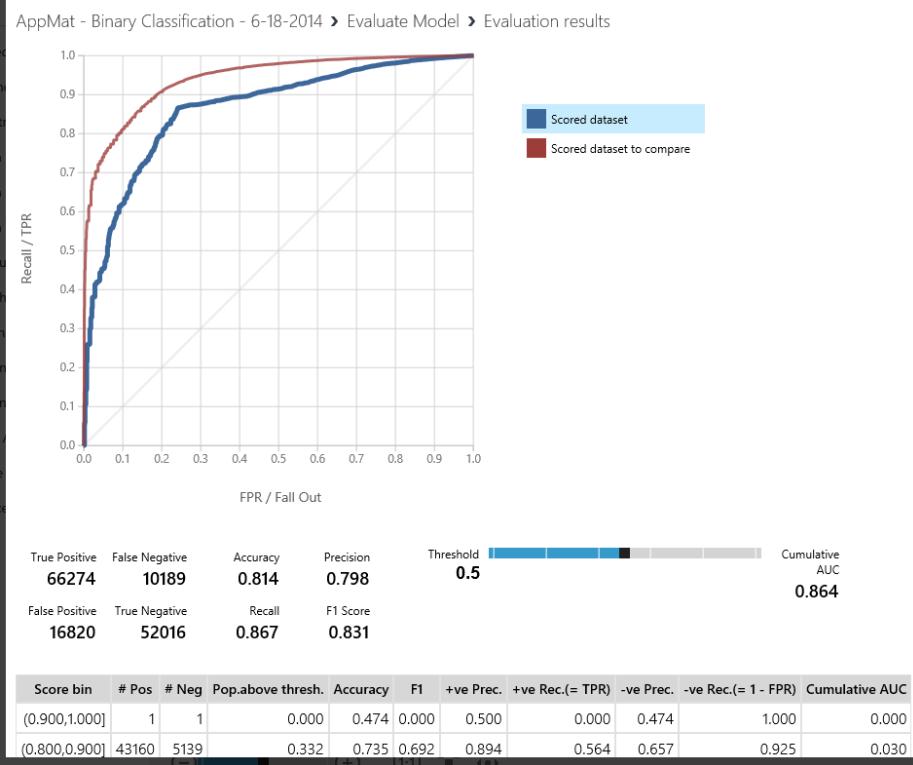
Split data to training and testing set

Check how model work on testing set



# Model Experiment and Evaluate

## Classification



- Accuracy =  $\frac{TP + TN}{P+N}$
- Precision =  $\frac{TP}{TP+FP}$
- Recall =  $\frac{TP}{TP+FN} = \frac{TP}{P}$
- $$F_\beta = \frac{(1+\beta^2) \cdot precision \cdot recall}{\beta^2 \cdot precision + recall}$$

## Regression

fy18-car-automobile-price > Evaluate Model >

### Metrics

Mean Absolute Error	3808.687462
Root Mean Squared Error	5114.95604
Relative Absolute Error	0.720506
Relative Squared Error	0.643097
Coefficient of Determination	0.356903

# Machine Learning on Azure

## Sophisticated pretrained models

To simplify solution development



## Popular frameworks

To build advanced deep learning solutions



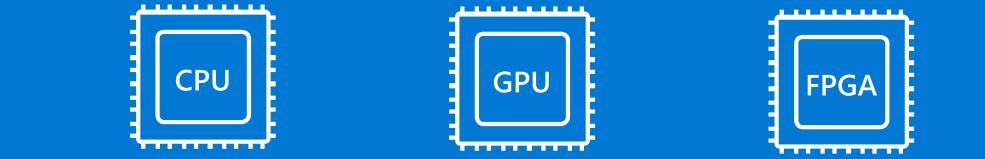
## Productive services

To empower data science and development teams



## Powerful infrastructure

To accelerate deep learning



## Flexible deployment

To deploy and manage models on intelligent cloud and edge



# Machine Learning & AI Portfolio

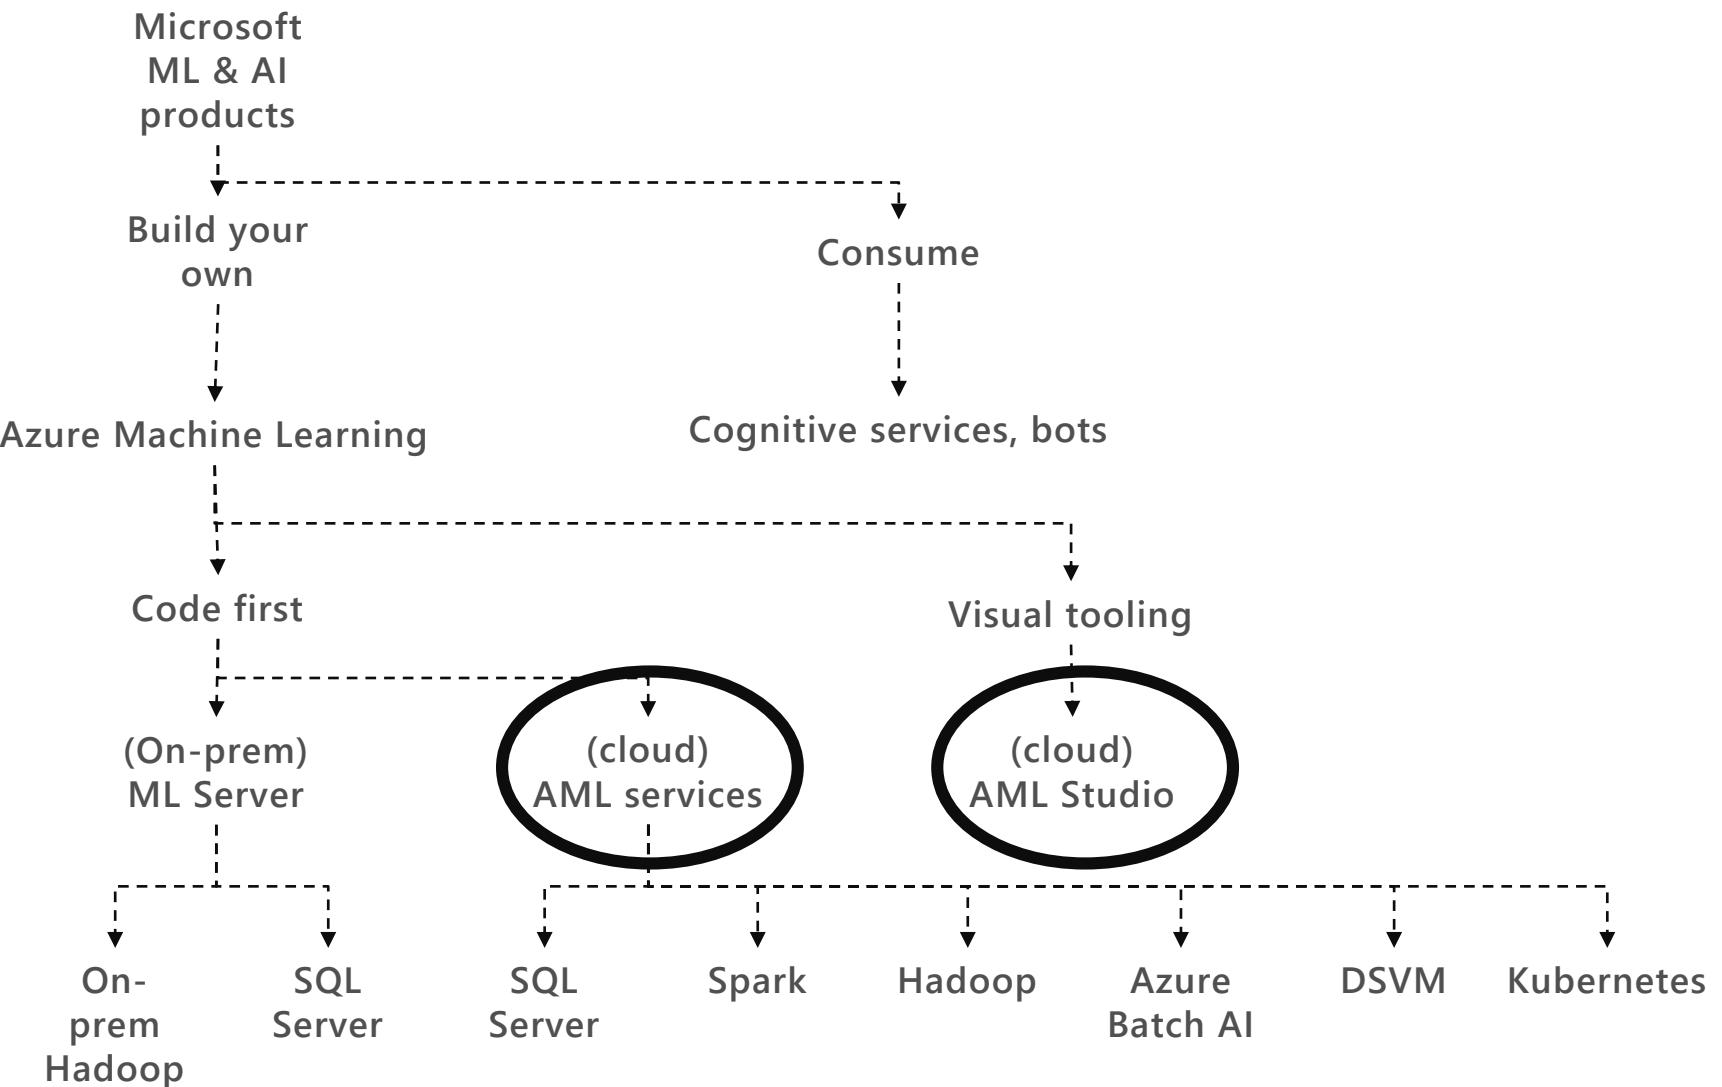
When to use what?

**Build your own or consume pre-trained models?**

**Which experience do you want?**

**Deployment target**

**What engine(s) do you want to use?**



# Learn – basic ML

<https://docs.microsoft.com/en-us/azure/machine-learning>

Watch/Read: [Data Science for beginners](#)

<https://docs.microsoft.com/en-us/azure/machine-learning/studio/data-science-for-beginners-the-5-questions-data-science-answers>

For more: Machine Learning & AI Foundations: Recommendations on [LinkedIn](#) (<https://www.linkedin.com/learning/machine-learning-ai-foundations-recommendations?u=3322>)

Python basics: [Python Essential Training](#) (<https://www.linkedin.com/learning/python-essential-training-2?u=3322> )

[Python for Machine Learning using jupyter notebooks](#) (<https://www.linkedin.com/learning/python-for-data-science-essential-training?u=3322> )

# Materials – step 2

<https://msropendata.com/> - many datasets from different industries!

<http://www.misjaazure.pl>

<https://gallery.azure.ai/>

<https://aischool.microsoft.com/en-us/home>

<https://www.microsoft.com/en-us/research/lab/microsoft-research-ai/>

<https://github.com/MSRConnections/Azure-training-course>

<https://www.ailab.microsoft.com/experiments>

# Demo

How to start with Machine Learning Studio?

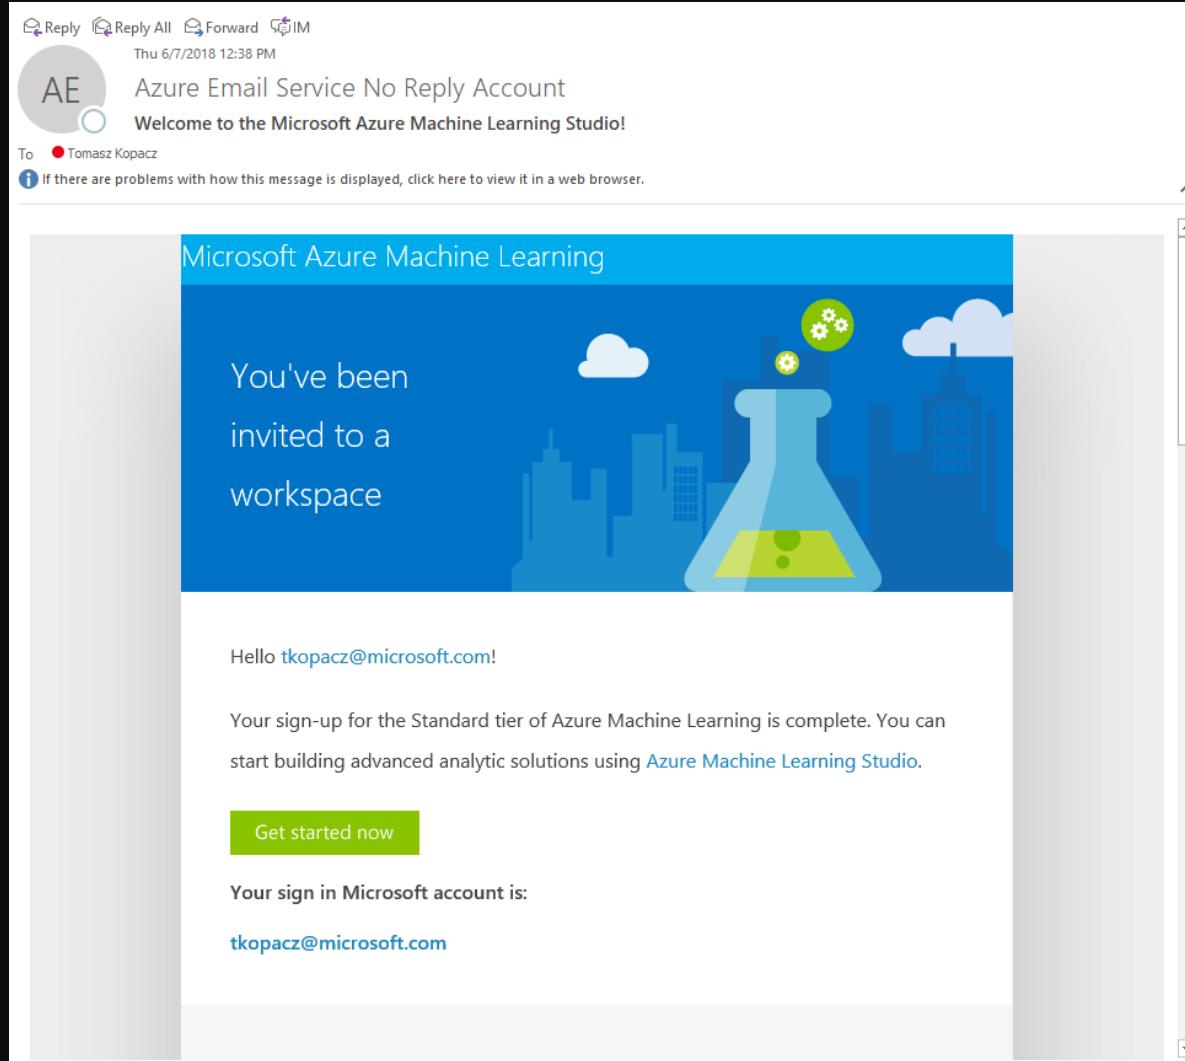
# Create Machine Learning Studio Workspace

<http://portal.azure.com>

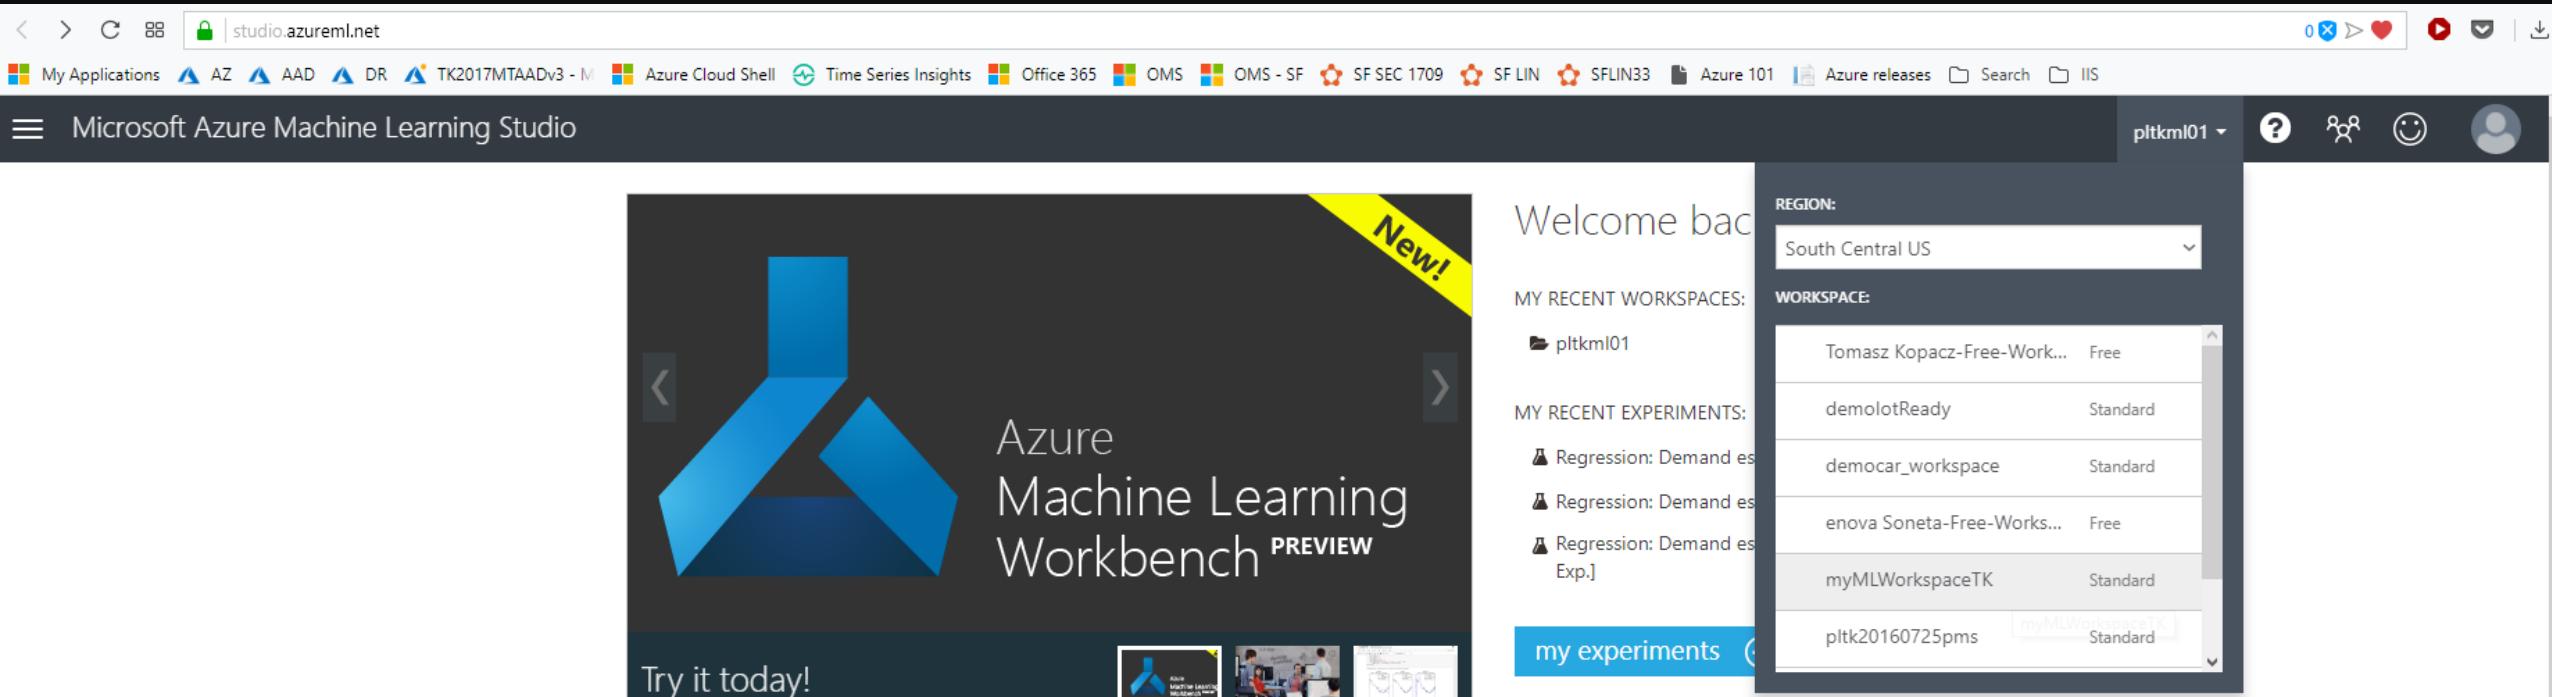
Create a Resource  
Machine Learning Studio Workspace

The screenshot shows the Microsoft Azure portal interface. On the left, the sidebar includes 'Create a resource', 'All services', 'Dashboard', 'Resource groups', 'App Services', 'Function Apps', 'SQL databases', 'Azure Cosmos DB', 'Virtual machines', 'Load balancers', 'Storage accounts', 'Virtual networks', 'Azure Active Directory', 'Monitor', 'Advisor', 'Security Center', 'Cost Management + Billing', and 'Help + support'. The main area shows a search bar with 'Machine Learning Studio Workspace' and a 'Popular' section with options like Windows Server 2016 VM, Ubuntu Server 17.10 VM, Web App, SQL Database, Serverless Function App, Cosmos DB, Kubernetes Service, DevOps Project, Storage Account, and Software as a service (SaaS). A central modal window titled 'Machine Learning Studio Workspace' provides a brief description and a 'Save for later' button. Below it, a preview window shows the 'Production workspace' settings. To the right, a large form for creating the workspace is displayed, with fields for 'Workspace name' (set to 'myMLWorkspaceTK'), 'Subscription' (set to 'HACKATON01'), 'Resource group' (radio button selected for 'Create new' and value 'rgML'), 'Location' (set to 'South Central US'), 'Storage account' (radio button selected for 'Create new' and value 'mymlworkspacektstorage'), 'Workspace pricing tier' (set to 'Standard'), 'Web service plan' (radio button selected for 'Create new' and value 'myMLWorkspaceTKPlan'), and 'Web service plan pricing tier' (set to 'Dev/Test Standard'). At the bottom of the form are 'Create' and 'Automation options' buttons.

# You should receive mail like this



# Go & login to studio.azure.ml



The screenshot shows the Microsoft Azure Machine Learning Studio homepage. The URL in the browser bar is `studio.azureml.net`. The top navigation bar includes links for My Applications, AZ, AAD, DR, TK2017MTAADv3 - M, Azure Cloud Shell, Time Series Insights, Office 365, OMS, OMS - SF, SF SEC 1709, SF LIN, SF LIN33, Azure 101, Azure releases, Search, IIS, and several other internal Azure services.

The main content area features a large graphic for the "Azure Machine Learning Workbench PREVIEW". It includes a "New!" badge, a "Try it today!" button, and three small preview images. To the right, a "Welcome back" message is displayed, along with dropdown menus for "REGION: South Central US" and "WORKSPACE". The workspace dropdown lists several entries:

Workspace Name	Plan
Tomasz Kopacz-Free-Work...	Free
demolotReady	Standard
democar_workspace	Standard
enova Soneta-Free-Works...	Free
myMLWorkspaceTK	Standard
pltk20160725pms	Standard

# Sample UI

Microsoft Azure Machine Learning Studio

ptkml01 ? ☺ 🔍

Regression: Demand estimation

Finished running ✓

Properties Project

Experiment Properties

START TIME 6/7/2018 1...  
END TIME 6/7/2018 1...  
STATUS CODE Finished  
STATUS DETAILS None

Summary

This experiment demonstrates demand estimation using regression with UCI bike rental data.

Description

Enter the detailed description for your experiment.

Original Experiment Documentation

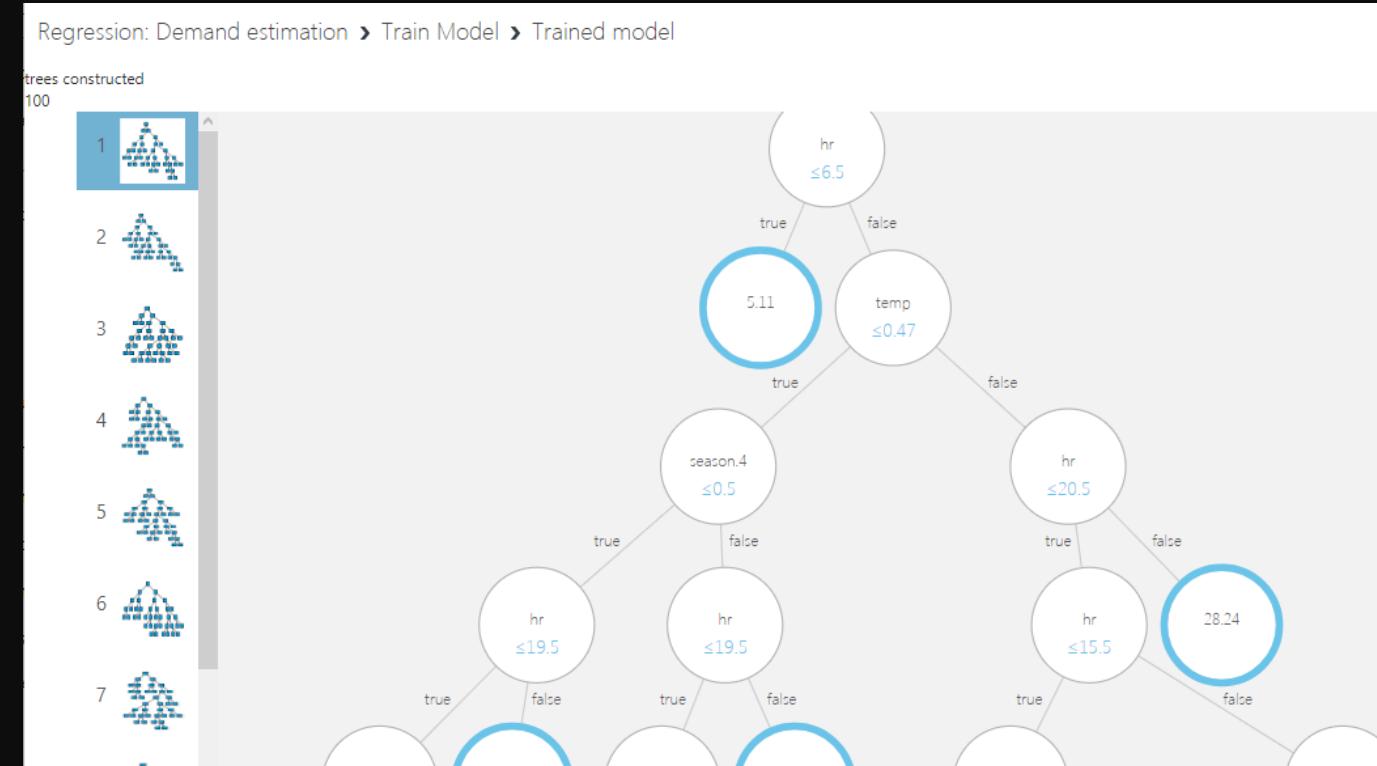
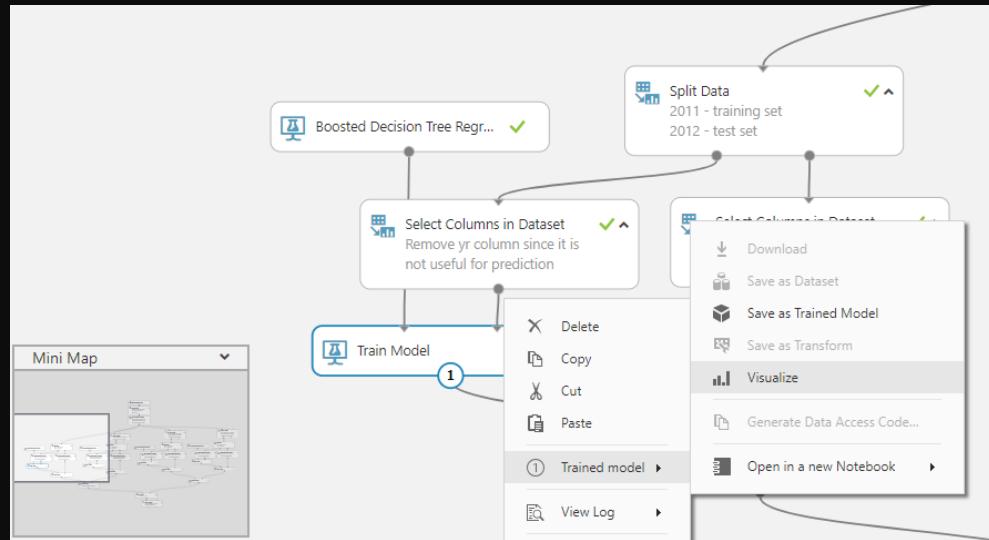
Quick Help

Mini Map

NEW RUN HISTORY SAVE DISCARD CHANGES RUN SET UP WEB SERVICE GALLERY

The screenshot shows a Microsoft Azure Machine Learning Studio experiment titled "Regression: Demand estimation". The experiment has completed successfully, indicated by the "Finished running" status and a green checkmark. The interface includes a left sidebar with various tools like Saved Datasets, Trained Models, and Data Input and Output, and a right sidebar with Experiment Properties, Summary, Description, and Original Experiment Documentation sections. The main area displays a complex workflow graph with nodes such as "Bike Rental UCI dataset", "Edit Metadata", "Selected Columns in Dataset", "Execute R Script", "Split Data", "Boosted Decision Tree Regr.", "Train Model", "Score Model", "Evaluate Model", "Add Rows", "Execute R Script", and "Selected Columns in Dataset". A "Mini Map" at the bottom provides a 3D view of the entire experiment structure. The bottom navigation bar includes icons for NEW, RUN HISTORY, SAVE, DISCARD CHANGES, RUN, SET UP WEB SERVICE, and GALLERY.

# See the model



# Sign

https://signup.azure.com/signup?offer=MS-AZR-0063P

Microsoft Azure tkopacz@tomaszkopacz.com Sign out

## Visual Studio Enterprise

Enjoy monthly credits and lower rates. Use MSDN software for development and test at no additional charge.



### 1 Agreement

I agree to the [subscription agreement](#), [offer details](#), and [privacy statement](#)

I would like information, tips, and offers about Azure, including Azure Newsletter and pricing updates, and other Microsoft products and services. [privacy statement](#)

**Sign up**

# Sample UI

Microsoft Azure Machine Learning Studio

ptkml01 ? ☺ 🔍

Regression: Demand estimation

Finished running ✓

Properties Project

Experiment Properties

START TIME 6/7/2018 1...  
END TIME 6/7/2018 1...  
STATUS CODE Finished  
STATUS DETAILS None

Summary

This experiment demonstrates demand estimation using regression with UCI bike rental data.

Description

Enter the detailed description for your experiment.

Original Experiment Documentation

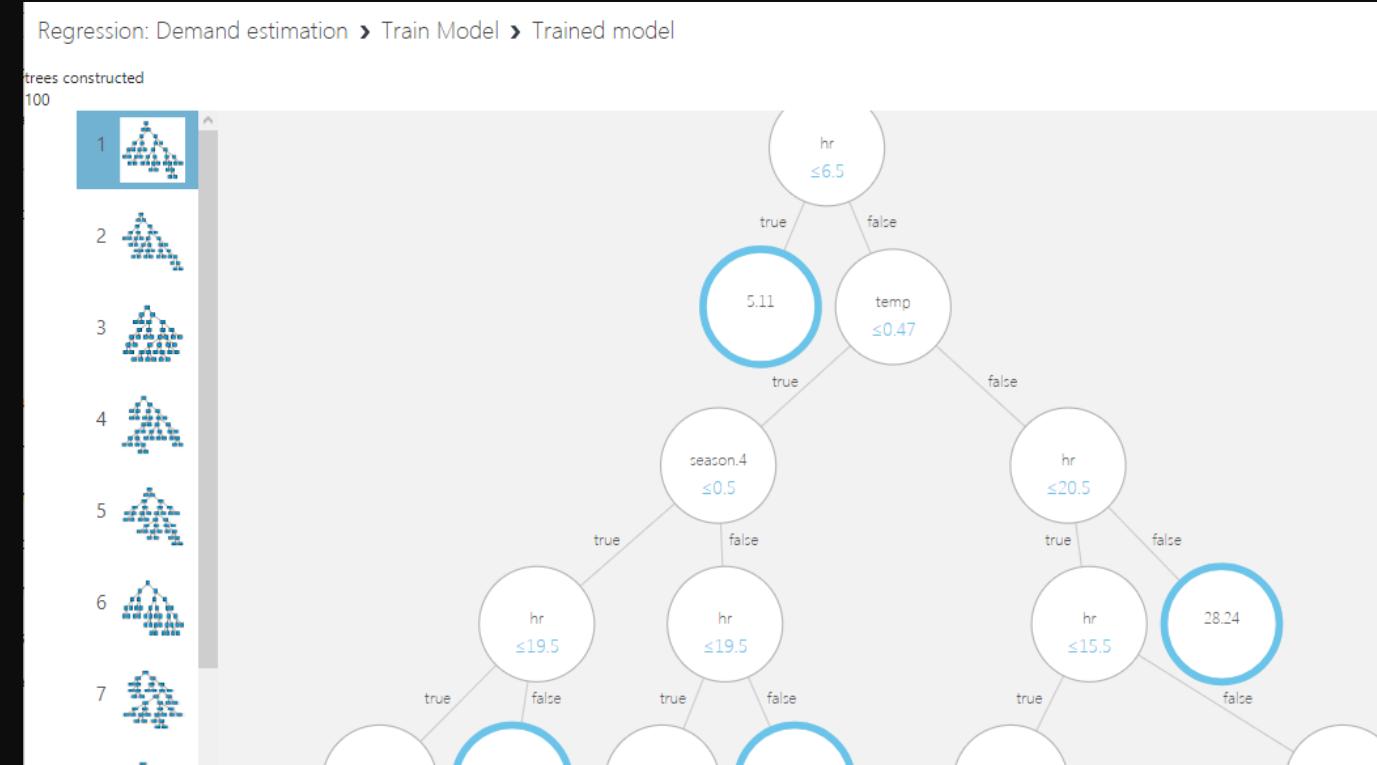
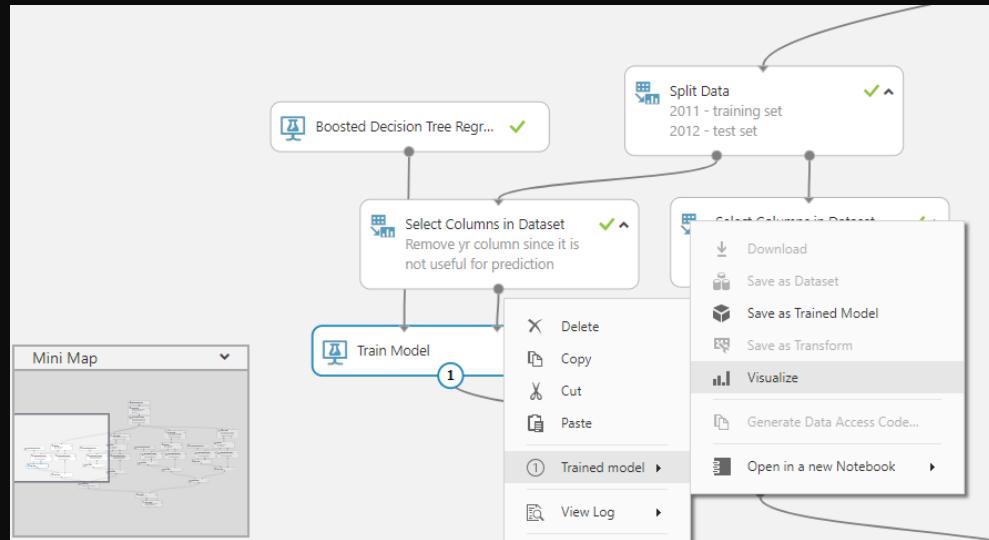
Quick Help

Mini Map

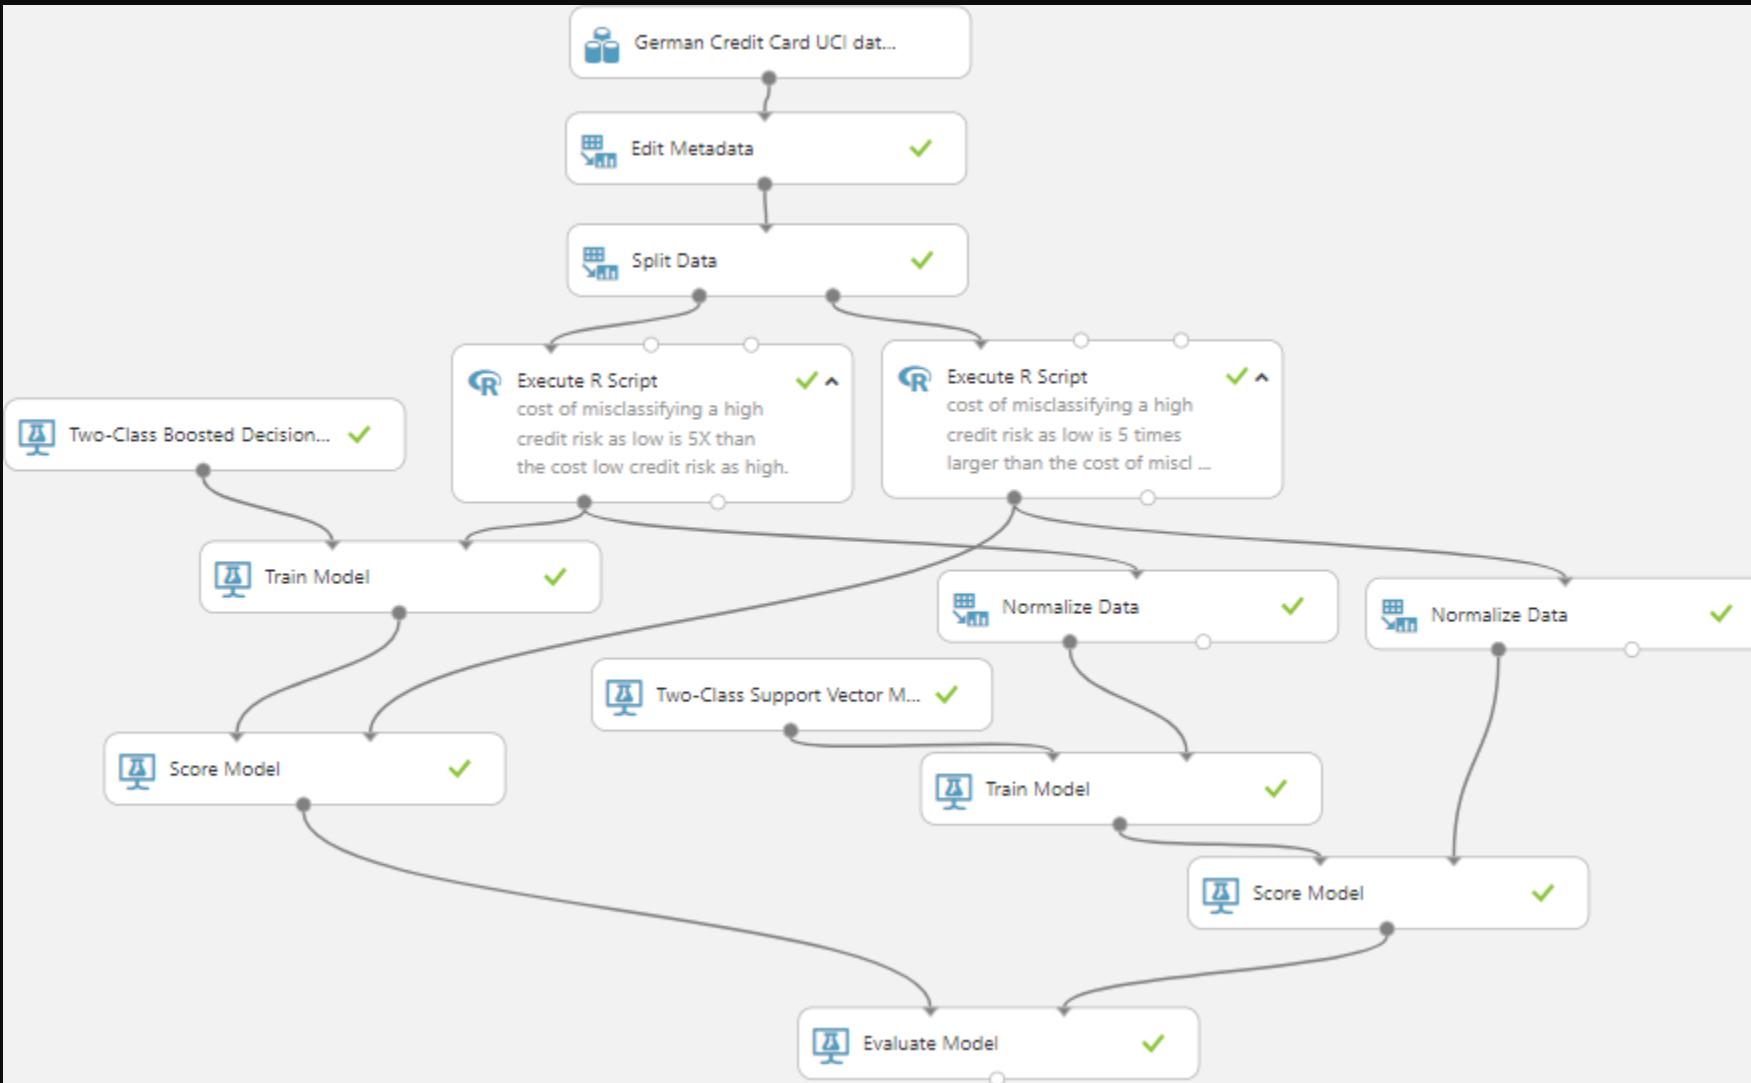
NEW RUN HISTORY SAVE DISCARD CHANGES RUN SET UP WEB SERVICE GALLERY

The screenshot shows a Microsoft Azure Machine Learning Studio experiment titled "Regression: Demand estimation". The experiment has completed successfully, indicated by the "Finished running" status and a green checkmark. The interface includes a left sidebar with various tools like Saved Datasets, Trained Models, and Data Input and Output, and a right sidebar with Experiment Properties, Summary, Description, and Original Experiment Documentation sections. The main workspace displays a complex workflow graph. At the top, a "Bike Rental UCI dataset" is processed through "Edit Metadata" and "Selected Columns in Dataset" steps. This leads to multiple parallel paths for data splitting and model training. There are four main splits: one path leads to "Boosted Decision Tree Regr." and "Score Model"; another path leads to "Execute R Script" and "Split Data"; a third path leads to "Select Columns in Dataset" and "Train Model"; and a fourth path leads to "Evaluate R Script" and "Split Data". These steps are interconnected with various "Select Columns in Dataset" and "Train Model" blocks, creating a dense network of operations. A "Mini Map" at the bottom provides a visual overview of the entire workflow structure.

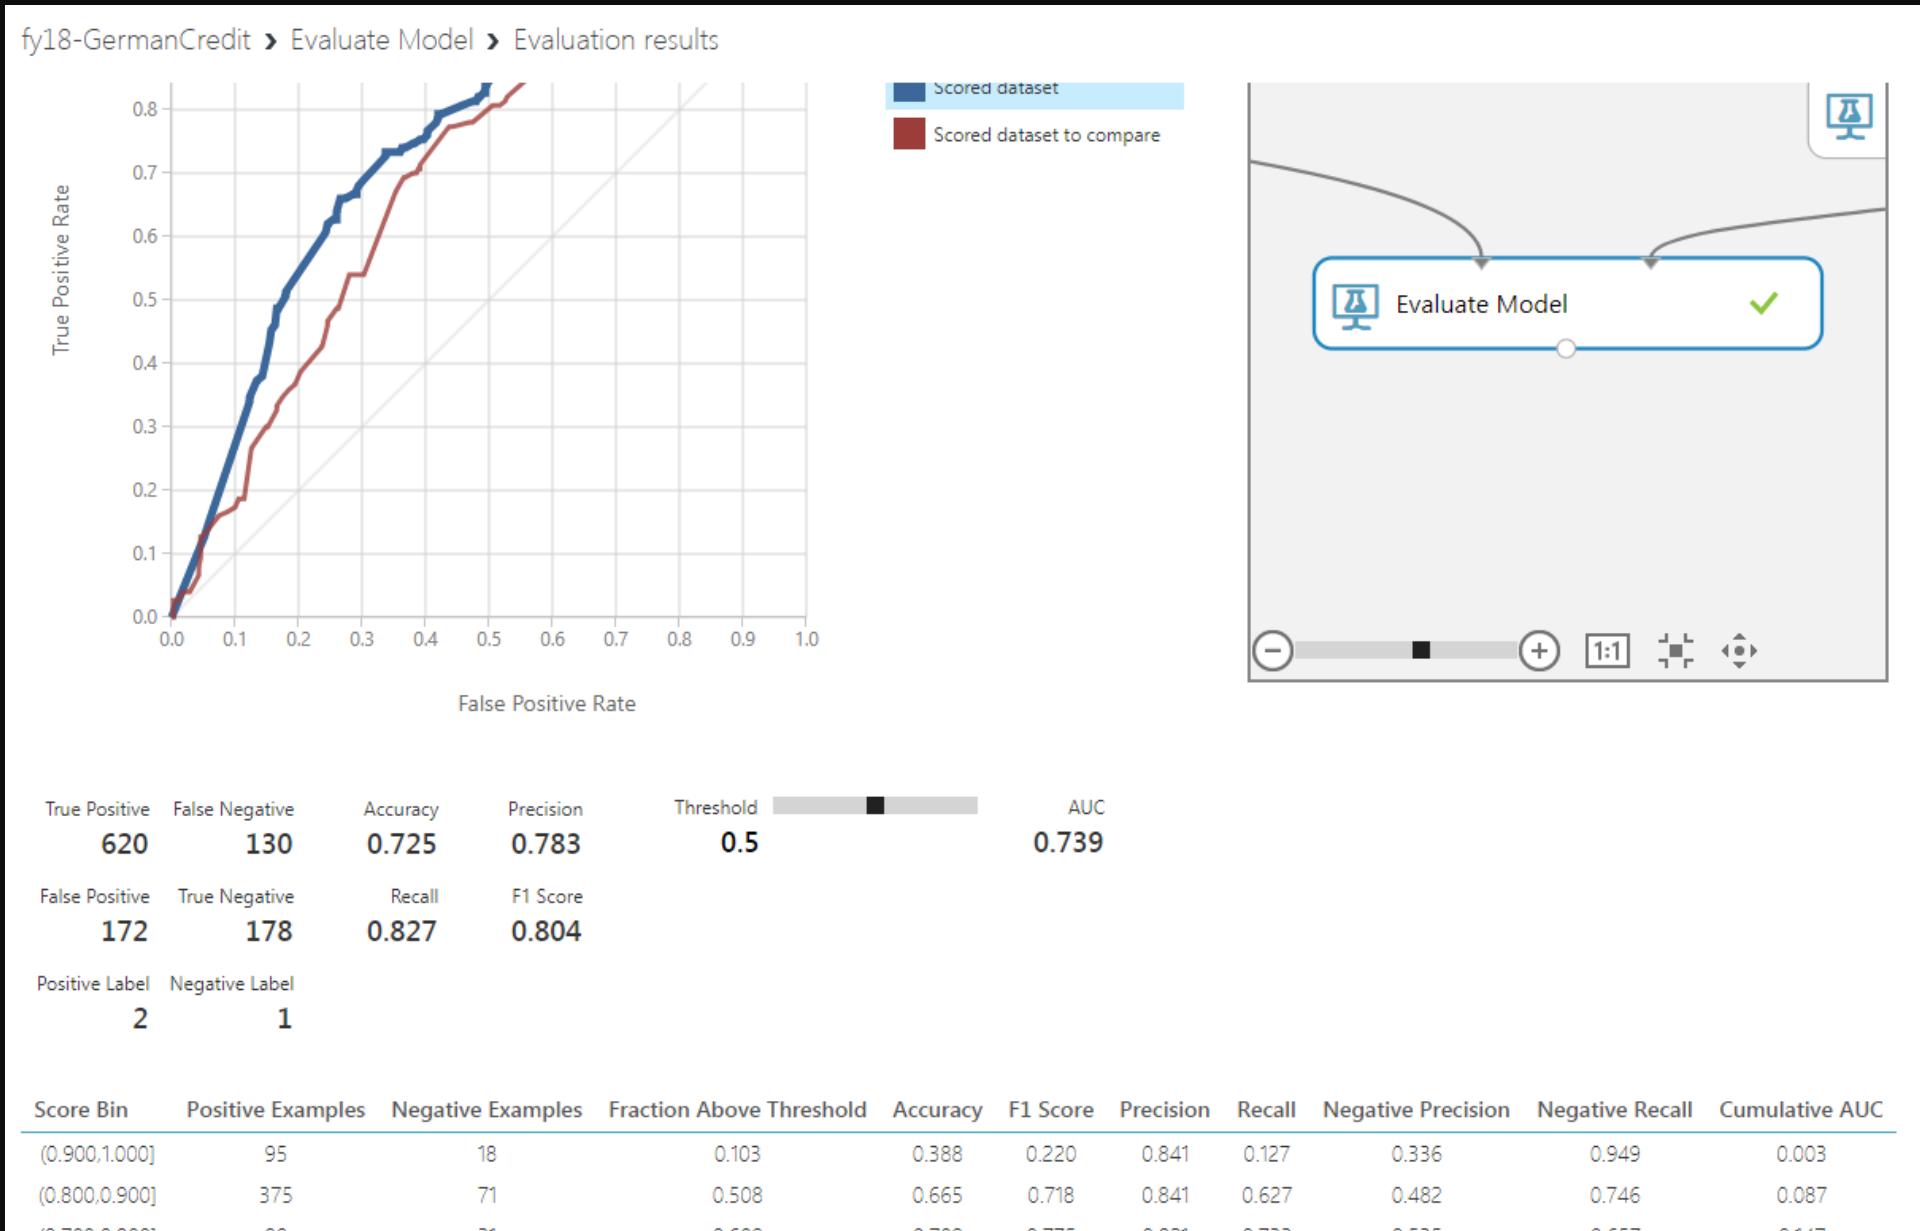
# See the model



# Final



# Results



# Publish better model as web service

Microsoft Azure Machine Learning Studio

Training experiment Predictive experiment

fy18-GermanCredit - to publish [Predictive Exp.]

Web service input

German Credit Card UCI dat...

Edit Metadata

Execute R Script  
cost of misclassifying a high credit risk as low is 5 times larger than the cost of miscl ...

fy18-GermanCredit - to pub...

Score Model

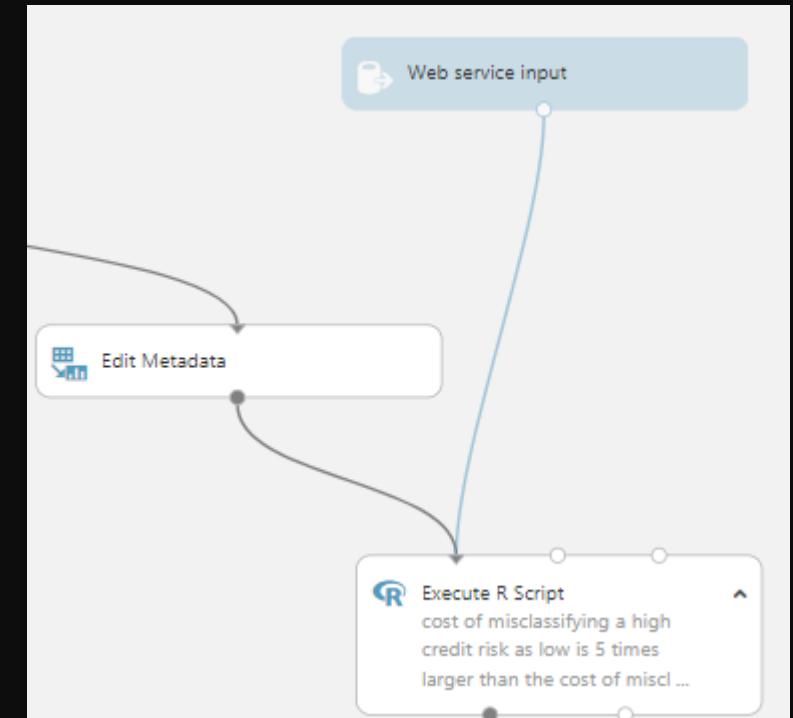
Score Model.  
Score a trained classification or regression model

Web service output

Creating predictive experiment

NEW RUN HISTORY SAVE SAVE AS DISCARD CHANGES RUN DEPLOY WEB SERVICE PUBLISH TO GALLERY

```
graph TD; WSInput[Web service input] --> GCUCI[German Credit Card UCI dat...]; GCUCI --> EM[Edit Metadata]; EM --> ERS[Execute R Script]; ERS --> FGY[fy18-GermanCredit - to pub...]; FGY --> SM[Score Model]; SM --> WSOutput[Web service output]
```



# Microsoft Cognitive Toolkit (CNTK)

1st-class on Linux and Windows and Docker

Rich API support

Mostly implemented in C++ (train and eval)

Low level + high level Python API

R and C# API for train and eval, UWP, Java and Spark support

Built-in readers for distributed learning

Keras backend support (Beta) (Keras = high-level neural networks API, for HUMANS)

Model compression (Fast binarized evaluation)

Latest version: v2.2 (Sep. 15, 2017)

Open Neural Network Exchange (ONNX)

(beta) Share model across frameworks.

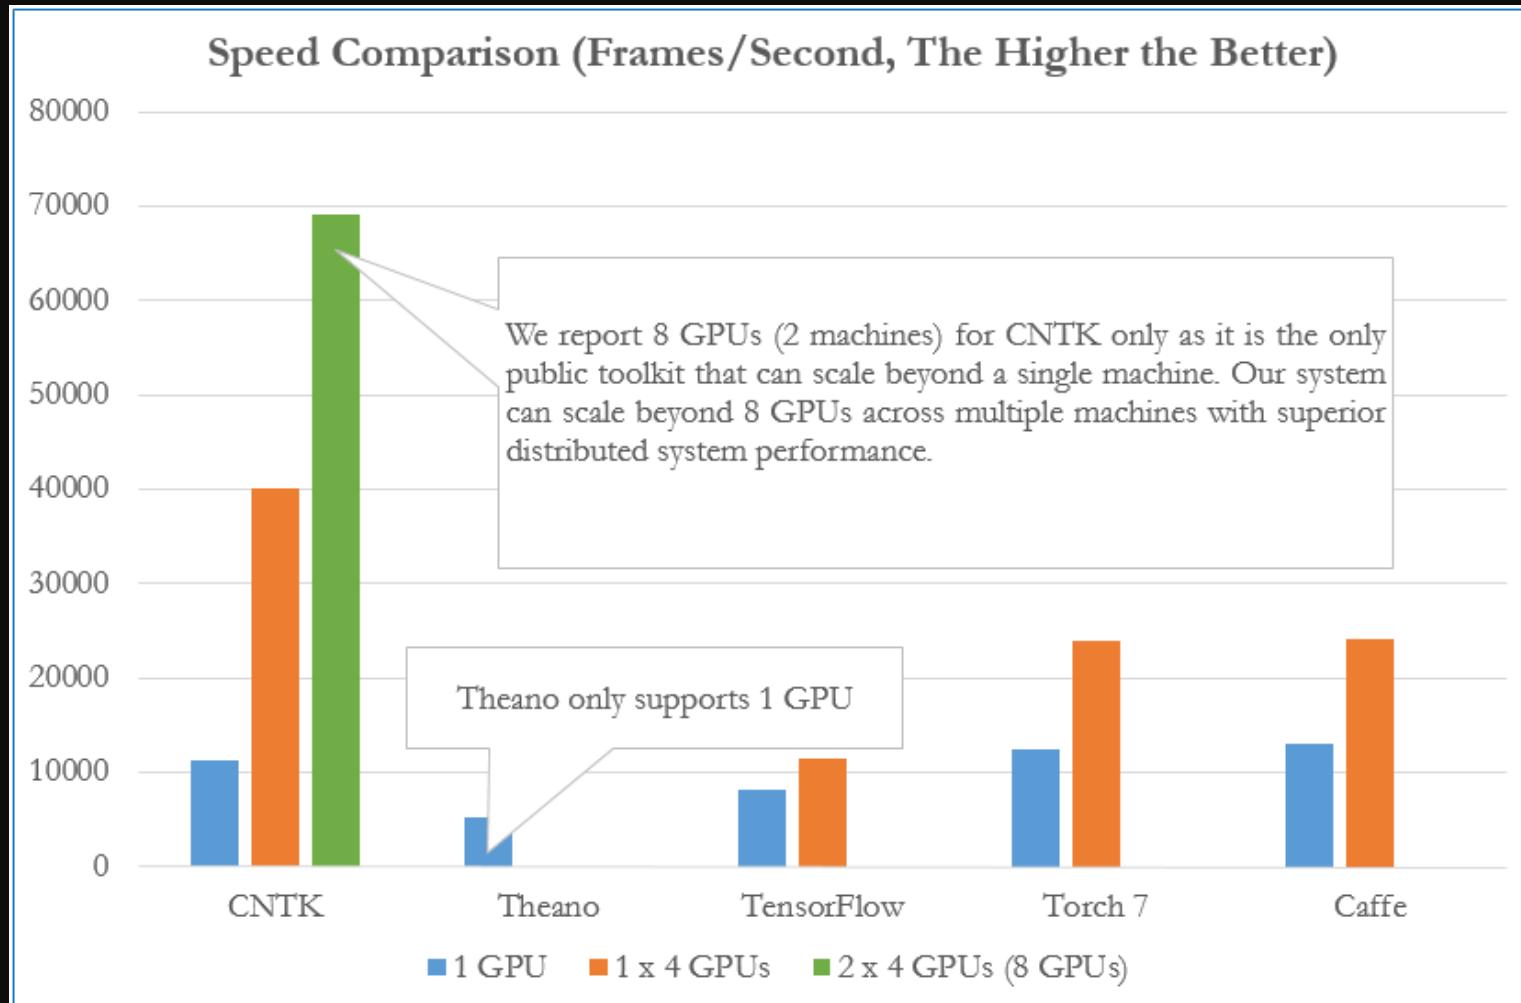
Calculate on CPU | GPU | FPGA | Silicon

Use / Train: CNTK / Caffe 2, PyTorch / Keras ....



# By the way – why CNTK ?

It can scale ACROSS MACHINES! [Docs here](#)



# SimpleNet.ipynb

jupyter SimpleNet Last Checkpoint: 14 hours ago (autosaved) Azure Notebooks My Libraries 2018Athens

File Edit View Insert Cell Kernel Data Widgets Help  
Run C Enter/Exit RISE Slideshow

```
progress_printer = ProgressPrinter(0)
trainer = C.Trainer(z, (ce, pe), [sgd(z.parameters, lr=lr_per_minibatch)], [progress_printer])

# Get minibatches of training data and perform model training
minibatch_size = 25
num_minibatches_to_train = 1024

aggregate_loss = 0.0
for i in range(num_minibatches_to_train):
    train_features, labels = generate_random_data(minibatch_size, inputs, outputs)
    # Specify the mapping of input variables in the model to actual minibatch data to be trained with
    trainer.train_minibatch({features : train_features, label : labels})
    sample_count = trainer.previous_minibatch_sample_count
    aggregate_loss += trainer.previous_minibatch_loss_average * sample_count

last_avg_error = aggregate_loss / trainer.total_number_of_samples_seen

test_features, test_labels = generate_random_data(minibatch_size, inputs, outputs)
avg_error = trainer.test_minibatch({features : test_features, label : test_labels})
print(' error rate on an unseen minibatch: {}'.format(avg_error))
return last_avg_error, avg_error
```

In [15]: ffnet()

average loss	since last	average metric	since last	examples
-----				
Learning rate per minibatch: 0.125				
0.917	0.917	0.56	0.56	25
0.788	0.723	0.52	0.5	75
0.734	0.693	0.463	0.42	175
0.651	0.579	0.395	0.335	375
0.579	0.511	0.277	0.168	775
0.5	0.423	0.202	0.129	1575
0.413	0.329	0.149	0.0963	3175
0.34	0.268	0.119	0.0903	6375
0.285	0.229	0.0988	0.0783	12775
0.25	0.216	0.0887	0.0786	25575
error rate on an unseen minibatch: 0.0				

Out[15]: (0.25024763371795417, 0.0)

The example above sets up a 2-layer fully connected deep neural network with 50 hidden dimensions per layer. We first setup two input variables, one for the input data and one for the labels. We then called the fully connected classifier network model function which simply sets up the required weights, biases, and activation functions for each layer.

We set two root nodes in the network: ce is the cross entropy which defined our model's loss function, and pe is the classification error. We set up a trainer object with the root nodes of the network and a learner. In this case we pass in the standard SGD learner with default parameters and a learning rate of 0.02.

# CNTK Toolkit – „HOL“

## Logistic Regression

A cancer hospital has provided data and wants us to determine if a patient has a fatal malignant cancer vs. a benign growth. To help classify each patient, we are given their age and the size of the tumor. Intuitively, one can imagine that younger patients and/or patient with small tumor size are less likely to have malignant cancer.

[https://cntk.ai/pythondocs/CNTK\\_101\\_LogisticRegression.html](https://cntk.ai/pythondocs/CNTK_101_LogisticRegression.html)

## LSTM Time Series (based on $\text{SIN}(x)$ )

Time Series Analysis – trying to predict „future“ value of  $f(x) = \sin(x)$  without knowing what is „sinus“ function

[https://cntk.ai/pythondocs/CNTK\\_106A\\_LSTM\\_Timeseries\\_with\\_Simulated\\_Data.html](https://cntk.ai/pythondocs/CNTK_106A_LSTM_Timeseries_with_Simulated_Data.html)

# Logistic Regression (malignant cancer)

Jupyter CNTK\_101\_LogisticRegression Last Checkpoint: 15 hours ago (autosaved) Azure Notebooks My Libraries 2018Athens

File Edit View Insert Cell Kernel Data Widgets Help Not Trusted Python 3.6 C

Visualization

It is desirable to visualize the results. In this example, the data is conveniently in two dimensions and can be plotted. For data with higher dimensions, visualization can be challenging. There are advanced dimensionality reduction techniques that allow for such visualizations [t-sne](#).

```
In [27]: # Model parameters
print(mydict['b'].value)

bias_vector = mydict['b'].value
weight_matrix = mydict['w'].value

# Plot the data
import matplotlib.pyplot as plt

# given this is a 2 class
colors = ['r' if l == 0 else 'b' for l in labels[:,0]]
plt.scatter(features[:,0], features[:,1], c=colors)
plt.plot([0, bias_vector[0]/weight_matrix[0][1]],
         [bias_vector[1]/weight_matrix[0][0], 0], c = 'g', lw = 3)
plt.xlabel("Scaled age (in yrs)")
plt.ylabel("Tumor size (in cm)")
plt.show()

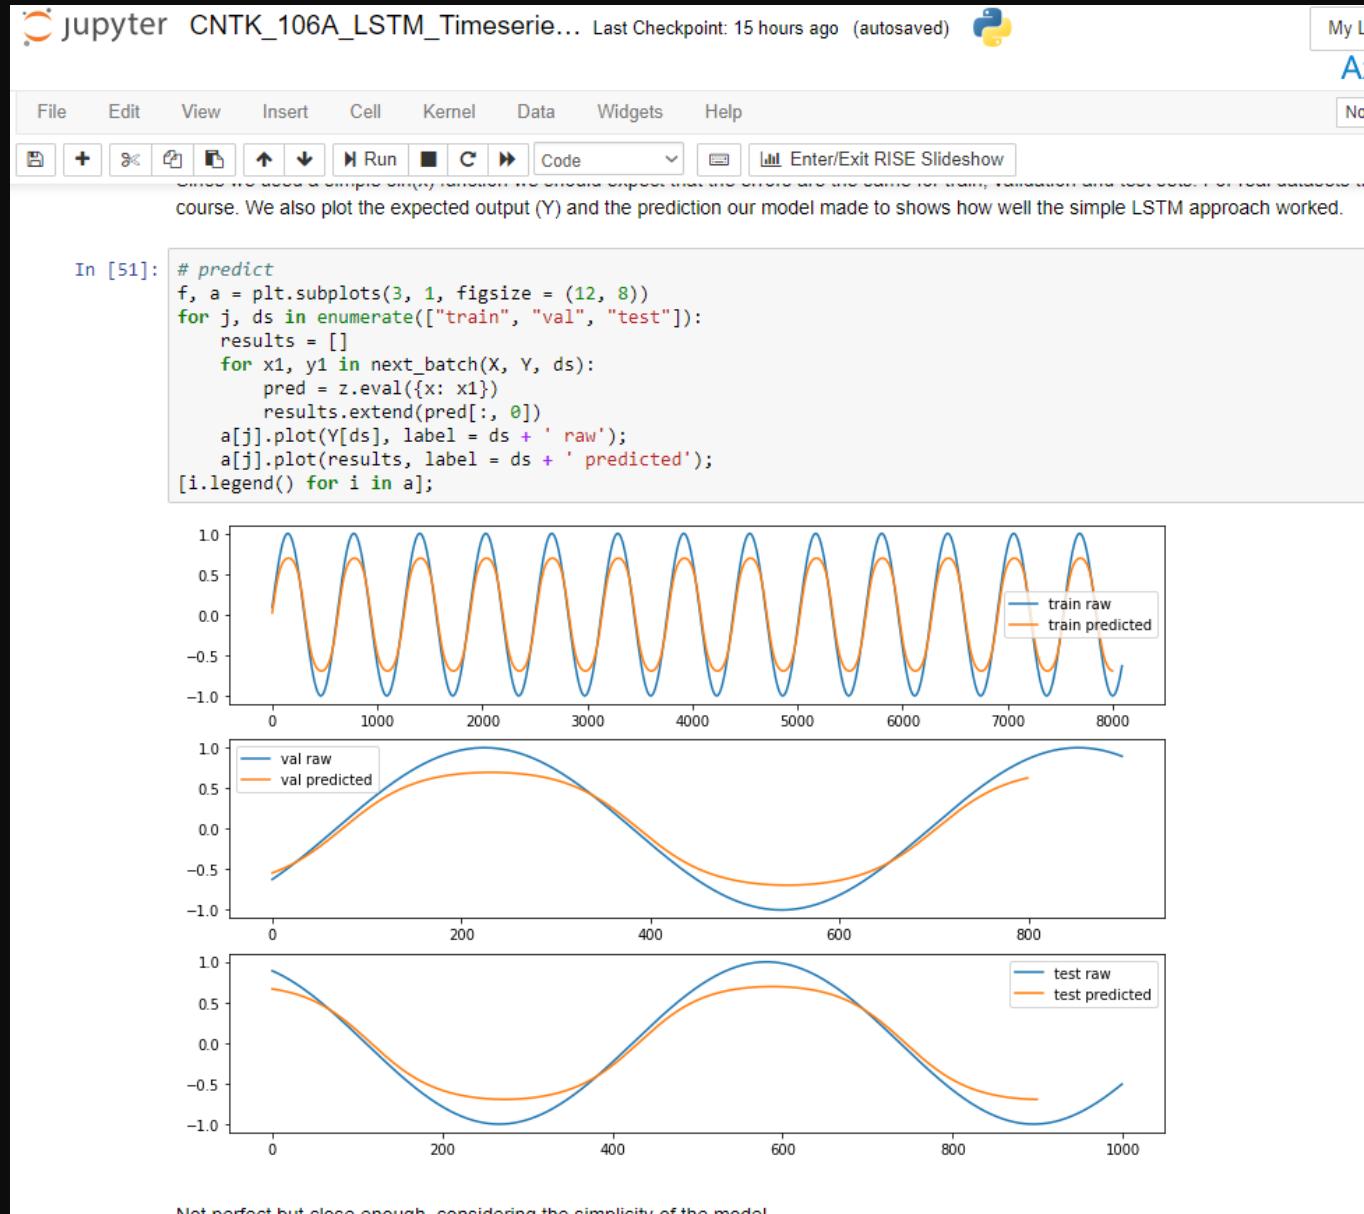
[ 7.991415 -7.991414]
```

A scatter plot showing tumor size (in cm) on the y-axis versus scaled age (in yrs) on the x-axis. The x-axis ranges from 0 to 10, and the y-axis ranges from 0 to 10. Red dots represent one class, and blue dots represent the other. A solid green line represents the decision boundary, which is a linear equation derived from the logistic regression model parameters.

Exploration Suggestions

- Try exploring how the classifier behaves with different data distributions - suggest changing the `minibatch_size` parameter from 25 to say 64. Why is the error increasing?
- Try exploring different activation functions
- Try exploring different learners
- You can explore training a [multiclass logistic regression](#) classifier.

# LSTM Time Series (based on SIN(x))



# And, remember



**Mat Velloso**  
@matvelloso

[Follow](#)



Half of the time when companies say they need "AI" what they really need is a SELECT clause with GROUP BY.

You're welcome.

# Demo

How to start with Machine Learning Services?

# Why Python? And Code

Customers have told us they love the convenience

Customers have told us they need:

Greater control over compute & data

More options for model deployment

Which frameworks? ALL OF THEM!

# Top 20 Python libraries for data science in 2018

[NumPy](#) (Commits: 17911, Contributors: 641)

[SciPy](#) (Commits: 19150, Contributors: 608)

[Pandas](#) (Commits: 17144, Contributors: 1165)

[StatsModels](#) (Commits: 10067, Contributors: 153)

[Matplotlib](#) (Commits: 25747, Contributors: 725)

[Seaborn](#) (Commits: 2044, Contributors: 83)

[Plotly](#) (Commits: 2906, Contributors: 48)

[Bokeh](#) (Commits: 16983, Contributors: 294)

[Pydot](#) (Commits: 169, Contributors: 12)

[Scikit-learn](#) (Commits: 22753, Contributors: 1084)

[XGBoost](#) / [LightGBM](#) / [CatBoost](#) (Commits: 3277 / 1083 / 1509,  
Contributors: 280 / 79 / 61)

[Eli5](#) (Commits: 922, Contributors: 6)

[TensorFlow](#) (Commits: 33339, Contributors: 1469)

[PyTorch](#) (Commits: 11306, Contributors: 635)

[Keras](#) (Commits: 4539, Contributors: 671)

[Dist-keras](#) / [elephas](#) / [spark-deep-learning](#)(Commits: 1125 / 170 / 67,  
Contributors: 5 / 13 / 11)

[NLTK](#) (Commits: 13041, Contributors: 236)

[SpaCy](#) (Commits: 8623, Contributors: 215)

[Gensim](#) (Commits: 3603, Contributors: 273)

[Scrapy](#) (Commits: 6625, Contributors: 281)

[Microsoft Cognitive Toolkit \(CNTK\)](#)

[Caffe2](#)

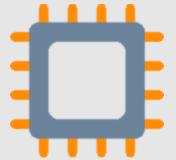
<100 more libraries ☺>

# Azure Machine Learning Service

# Team Workspace - Logical



Workspace



compute



models



experiments



images



data stores



deployment

# Team Workspace - physical



workspace



storage



container  
registry



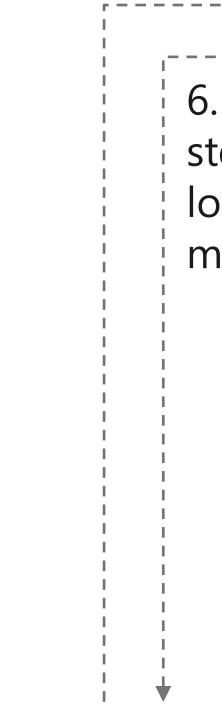
key vault



application  
insights

# Workflow

1. Snapshot folder and send to experiment



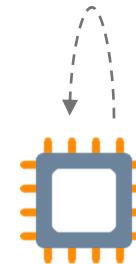
6. Stream  
stdout,  
logs,  
metrics

2. create docker image



Docker Image

5. Launch script



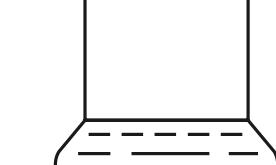
7. Copy over outputs

3. Deploy docker  
and snapshot to  
compute



4. Mount datastore  
to compute

Data Store



My Computer

Azure ML  
Workspace

## Prepare



Prepare  
Data

## Experiment



Build model  
(your favorite  
IDE)



Train & Test  
Model



Register and  
Manage Model



Build  
Image



Deploy  
Service  
Monitor  
Model

## Deploy

# Types of Compute Target

## Virtual Machine

- New in Azure | Existing in Azure
- Anywhere (on prem as well!)

## Azure Batch AI

- Best choice regarding price/performance
- + Low priority instances

## Azure Kubernetes Service (AKS)

- Maybe for CPU bound? Many parallel calculations? Existing „free“ cluster?

### Add Compute

Compute name

Compute type  ▼

Create new  Use existing

VIRTUAL MACHINE DETAILS

 Please note that Azure Machine Learning service only supports virtual machines running Ubuntu ↗

Address

SSH Port

User name

Authentication type  Password  SSH Key

Password

# VM on Azure

Type	Sizes	Description
General purpose	B, Dsv3, Dv3, DSv2, Dv2, Av2, DC	Balanced CPU-to-memory ratio. Ideal for testing and development, small to medium databases, and low to medium traffic web servers.
Compute optimized	Fsv2, Fs, F	High CPU-to-memory ratio. Good for medium traffic web servers, network appliances, batch processes, and application servers.
Memory optimized	Esv3, Ev3, M, GS, G, DSv2, Dv2	High memory-to-CPU ratio. Great for relational database servers, medium to large caches, and in-memory analytics.
Storage optimized	Ls 	High disk throughput and IO. Ideal for Big Data, SQL, and NoSQL databases.
GPU	NV, NVv2, NC, NCv2, NCv3, ND	Specialized virtual machines targeted for heavy graphic rendering and video editing, as well as model training and inferencing (ND) with deep learning. Available with single or multiple GPUs.
High performance compute	H	Our fastest and most powerful CPU virtual machines with optional high-throughput network interfaces (RDMA).

# NC\_v2 – Pascal Generation GPU

- Pascal generation GPU instances - NVIDIA Tesla P100 GPUs
- Targeted for accelerating machine training jobs and HPC
- Specs:
  - FP64 – 4.7 TFLOPS of double precision floating point performance
  - FP32 – 9.3 TFLOPS of single precision performance
  - FP16 – 18.7 TFLOPS of half-precision performance
  - GPU Memory 16 GB

	NC6s_v2	NC12s_v2	NC24s_v2	NC24rs_v2
Cores (Broadwell 2.6Ghz)	6	12	24	24
GPU	1 x P100	2 x P100	4 x P100	4 x P100
Memory	112 GB	224 GB	448 GB	448 GB
Local Disk	~700 GB SSD	~1.4 TB SSD	~3 TB SSD	~3 TB SSD
Network	Azure Network	Azure Network	Azure Network	Azure Network + InfiniBand



# Batch Services

## Azure Batch Services

- Cloud-scale job scheduling and compute management
- Linux or Windows, any VM size
- [BatchLab](#) – alternative UI to manage

Flow:

- Upload data files and application files
- Create Pool of compute nodes
- Create Job
- Add task(s) to Job, pointing into
- Start computation (send job)

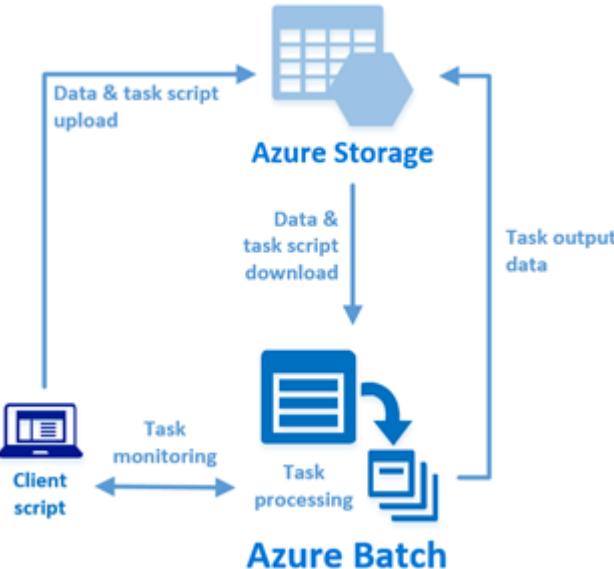
Two general types of task

- Parameter sweep (for each row, for set of number)
- Message Passing Interface – HPC style, multi-instance task

Work with clusters of GPUs

Autodesk, Chaos Group V-Ray, custom

[Shipyard](#) (containers; see also later)



ID	STATE	POOL
tkopacz20160122-093306job-0000000001	Active	mypool
tkopacz20160301-121612job-0000000001	Active	tikdem
f6a78dc7-3038-4a76-bc3d-55d193b20db9-20160831132154367job...	Completed	4ADEA

TASK	STATE	CREATED
task1	Completed	Jan 22 09:34:01
task2	Completed	Jan 22 09:34:01

# Azure Batch AI

To sum up: Batch Processing  
Pool – Job – Task (MPI, dependences etc)

Batch AI – specialized service

Fire and forget jobs, auto-scaling clusters, resource sharing

Batch or interactive use, VM images, containers, or setup script

Framework-specific configuration

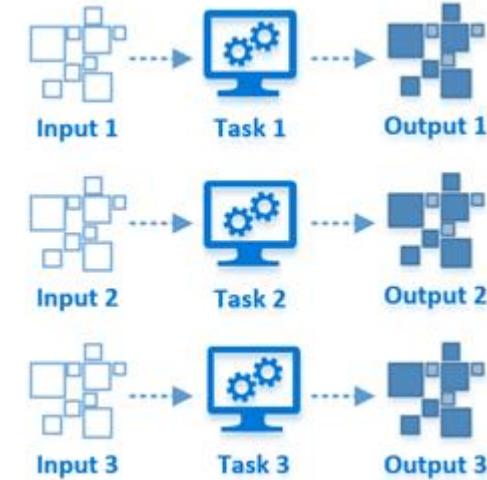
Stream output and return logs, Export trained models and configuration

Azure Files, Blob FUSE, NFS, or parallel file system

Lustre, Gluster, BeeGFS, ...

Azure CLI, Python, C#, Java, and REST API's, RBAC

Support for any Deep Learning or machine learning framework: [Microsoft Cognitive Toolkit \(CNTK\)](#), [TensorFlow](#), and [Chainer](#)



Standard and  
low priority VMs

# Train model

## Run (ACI, Batch AI)

```
from azureml.core.runconfig import RunConfiguration
from azureml.core.conda_dependencies import CondaDependencies

# create a new runconfig object
run_config = RunConfiguration()

# signal that you want to use ACI to execute script.
run_config.target = "containerinstance"

# ACI container group is only supported in certain regions, which can be different to
run_config.container_instance.region = 'eastus2'

# set the ACI CPU and Memory
run_config.container_instance.cpu_cores = 1
run_config.container_instance.memory_gb = 2

# enable Docker
run_config.environment.docker.enabled = True

# set Docker base image to the default CPU-based image
run_config.environment.docker.base_image = azureml.core.runconfig.DEFAULT_CPU_IMAGE

# use conda_dependencies.yml to create a conda environment in the Docker image for each
run_config.environment.python.user_managed_dependencies = False

# auto-prepare the Docker image when used for execution (if it is not already prepared)
run_config.auto_prepare_environment = True

# specify CondaDependencies obj
run_config.environment.python.conda_dependencies = CondaDependencies.create(conda_packages)

compute_config = BatchAiCompute.provisioning_configuration(vm_size="STANDARD_NC6", # GPU-based VM
                                                               #vm_priority='Lowpriority', # optional
                                                               autoscale_enabled=True,
                                                               cluster_min_nodes=0,
                                                               cluster_max_nodes=4)

# create the cluster
compute_target = ComputeTarget.create(ws, batchai_cluster_name, compute_config)
```

```
%time
from azureml.core.script_run_config import ScriptRunConfig

script_run_config = ScriptRunConfig(source_directory='./',
                                    script='train.py',
                                    run_config=run_config)

run = experiment.submit(script_run_config)

CPU times: user 344 ms, sys: 109 ms, total: 453 ms
Wall time: 5.2 s

# Show run details
run

%%time
# Shows output of the run on stdout.
run.wait_for_completion(show_output=True)
libstdc++-ng-8.2.0 | 2.9 MB | #####/##### | 78%
libstdc++-ng-8.2.0 | 2.9 MB | #####/##### | 100%

libedit-3.1 | 171 KB | | 0%
libedit-3.1 | 171 KB | #####/##### | 100%

mkl_fft-1.0.6 | 150 KB | | 0%
mkl_fft-1.0.6 | 150 KB | #####/##### | 100%
Downloading and Extracting Packages
Preparing transaction: ...working... done
Verifying transaction: ...working... done
Executing transaction: ...working... done

Collecting azureml-defaults==0.1.68 (from -r /azureml-setup/
    Downloading https://files.pythonhosted.org/packages/46/ad/
azur..._defaults-0.1.68-py2.py3-none-any.whl
Collecting azureml-core==0.1.68.* (from azureml-defaults==0
ine 1))
```

```
zureuser@pltkw3vmgpu38c8932019:~$ nvidia-smi
ed Oct 24 21:11:49 2018
+
+-----+-----+-----+-----+-----+-----+-----+-----+
| NVIDIA-SMI 396.44 | Driver Version: 396.44 |
+-----+-----+-----+-----+-----+-----+-----+-----+
| GPU  Name      Persistence-M| Bus-Id     Disp.A  | Volatile Uncorr. ECC |
| Fan  Temp  Perf  Pwr:Usage/Cap| Memory-Usage | GPU-Util  Compute M. |
|-----+-----+-----+-----+-----+-----+-----+-----+
| 0  Tesla K80        On          0000A20D:00:00.0 Off   |                0 |
| N/A   33C    P0    75W / 149W |    292MiB / 11441MiB |    14%     Default |
+-----+-----+-----+-----+-----+-----+-----+-----+
+
+-----+-----+-----+-----+-----+-----+-----+-----+
| Processes:                               GPU Memory |
| GPU     PID  Type  Process name        Usage     |
+-----+-----+-----+-----+-----+
| 0       16574  C    ...9c40c6245228f2ad32b89cddb31c/bin/python  281MiB |
+-----+-----+-----+-----+-----+-----+-----+-----+
zureuser@pltkw3vmgpu38c8932019:~$ █
```

# Demo

Training on:  
Local – end 2 end  
VM z ACI  
VM z GPU / Tensorboard  
Batch AI

# Register Model



# register model

```
# register model from the best run
model = best_run.register_model(model_name = 'mnist_tf_model', model_path = 'outputs/model')
```

Note the registered model automatically gets an auto-increasing version number.

```
print(model.name, model.id, model.version, sep = '\t')
mnist_tf_model  mnist_tf_model:2          2
```

Download the model folder locally and inspect the files in it.

```
import os
model.download('.')
os.listdir('./model/')

['mnist-tf.model.meta',
 'checkpoint',
 'mnist-tf.model.index',
 'mnist-tf.model.data-00000-of-00001']
```

# register model



Experiments Compute **Models** Images Deployments Activities

## mnist\_tf\_model

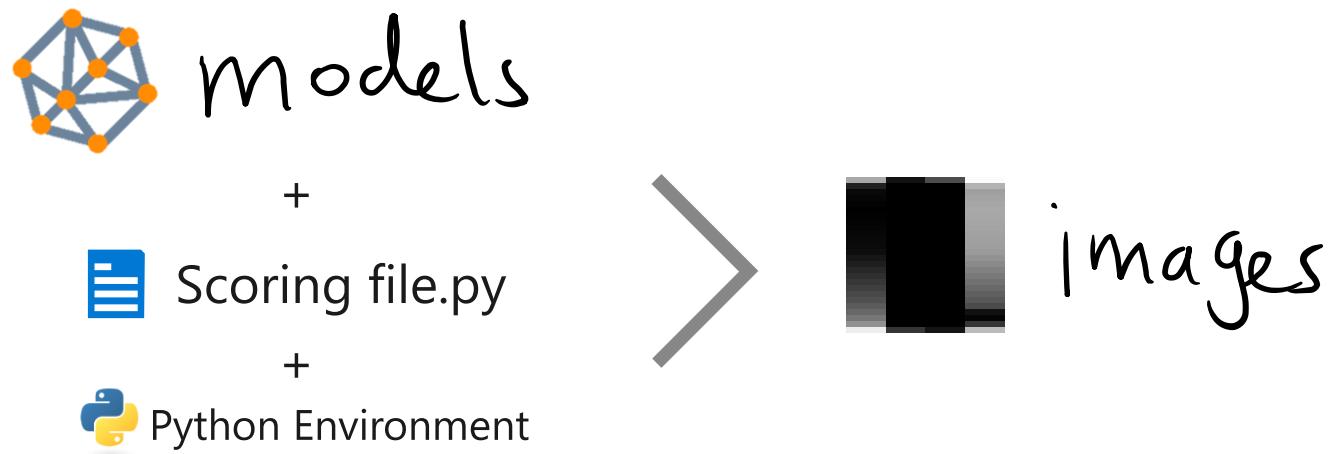
[Back to Models](#) [Delete](#)

[Details](#) [Deployments](#)

Attributes	
Version	1
ID	mnist_tf_model:1
Date Registered	09/16/2018, 5:22:29 PM UTC
Location	aml://asset/7e525521b9614308a993... <a href="#">Copy</a>
Description	
Tags	

# Create Image

# Create images



# Scoring File

```
1 import json
2 import numpy as np
3 import os
4 import tensorflow as tf
5
6 #from azureml.assets.persistence.persistence import get_model_path
7 from azureml.core.model import Model
8
9 def init():
10     global X, output, sess
11     tf.reset_default_graph()
12     # retreive the local path to the model using the model name
13     model_root = Model.get_model_path('mnist_tf_model')
14     saver = tf.train.import_meta_graph(os.path.join(model_root, 'mnist-tf.model.meta'))
15     X = tf.get_default_graph().get_tensor_by_name("network/X:0")
16     output = tf.get_default_graph().get_tensor_by_name("network/output/MatMul:0")
17
18     sess = tf.Session()
19     saver.restore(sess, os.path.join(model_root, 'mnist-tf.model'))
20
21 def run(raw_data):
22     data = np.array(json.loads(raw_data)[ 'data' ])
23     # make prediction
24     out = output.eval(session = sess, feed_dict = {X: data})
25     y_hat = np.argmax(out, axis = 1)
26     return json.dumps(y_hat.tolist())
```

# Environment File

```
# Conda environment specification. The dependencies defined in this file will
# be automatically provisioned for runs with userManagedDependencies=False.

# Details about the Conda environment file format:
# https://conda.io/docs/user-guide/tasks/manage-environments.html#create-env-file-manually

name: project_environment
dependencies:
    # The python interpreter version.
    # Currently Azure ML only supports 3.5.2 and later.
    - python=3.6.2

    - pip:
        # Required packages for AzureML execution, history, and data preparation.
        - azureml-defaults
        - tensorflow==1.9.0
```

# Create the Image

```
from azureml.core.image import Image, ContainerImage

image_config = ContainerImage.image_configuration(runtime= "python",
                                                 execution_script="score.py",
                                                 conda_file="myenv.yml")

image = Image.create(name = "mnist_image",
                     models = [model],                               # this is the model object
                     image_config = image_config,
                     workspace = ws)
```

# Images

Experiments Compute Models Images Deployments Activities

## tf-mnist

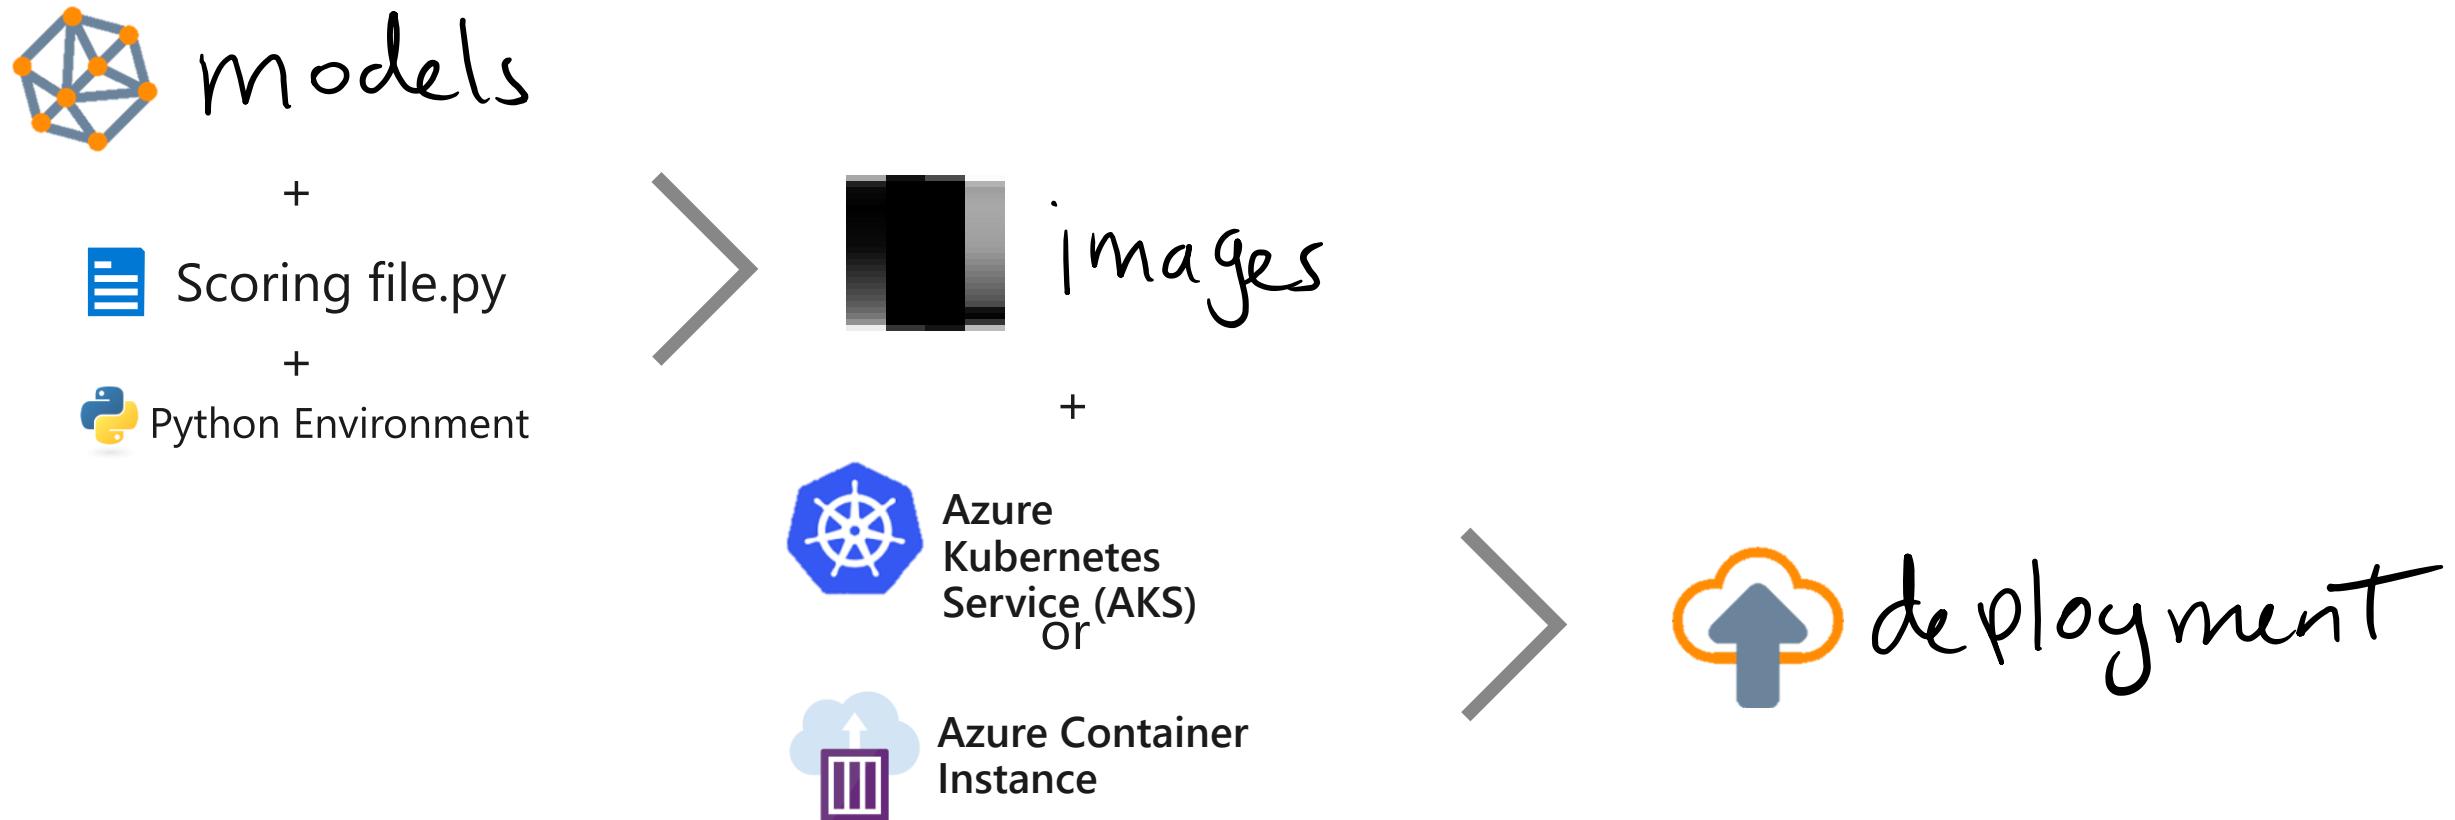
[← Back to Images](#) [+ Create Deployment](#)  [Delete](#)

[Details](#) [Models](#) [Deployments](#)

Attributes	
Description	
ID	tf-mnist:1
Date Registered	09/16/2018, 5:24:27 PM UTC
Version	1
Location	danielscacrancasv.azurecr.io/tf-mni... <a href="#"></a>
Environment	
Type	Docker
Status	Succeeded

# Deploy image

# Deploy image



# Deploy image

# Deploy image

Experiments Compute Models Images Deployments Activities

 Create deployment "aci-service-mnist": Success!

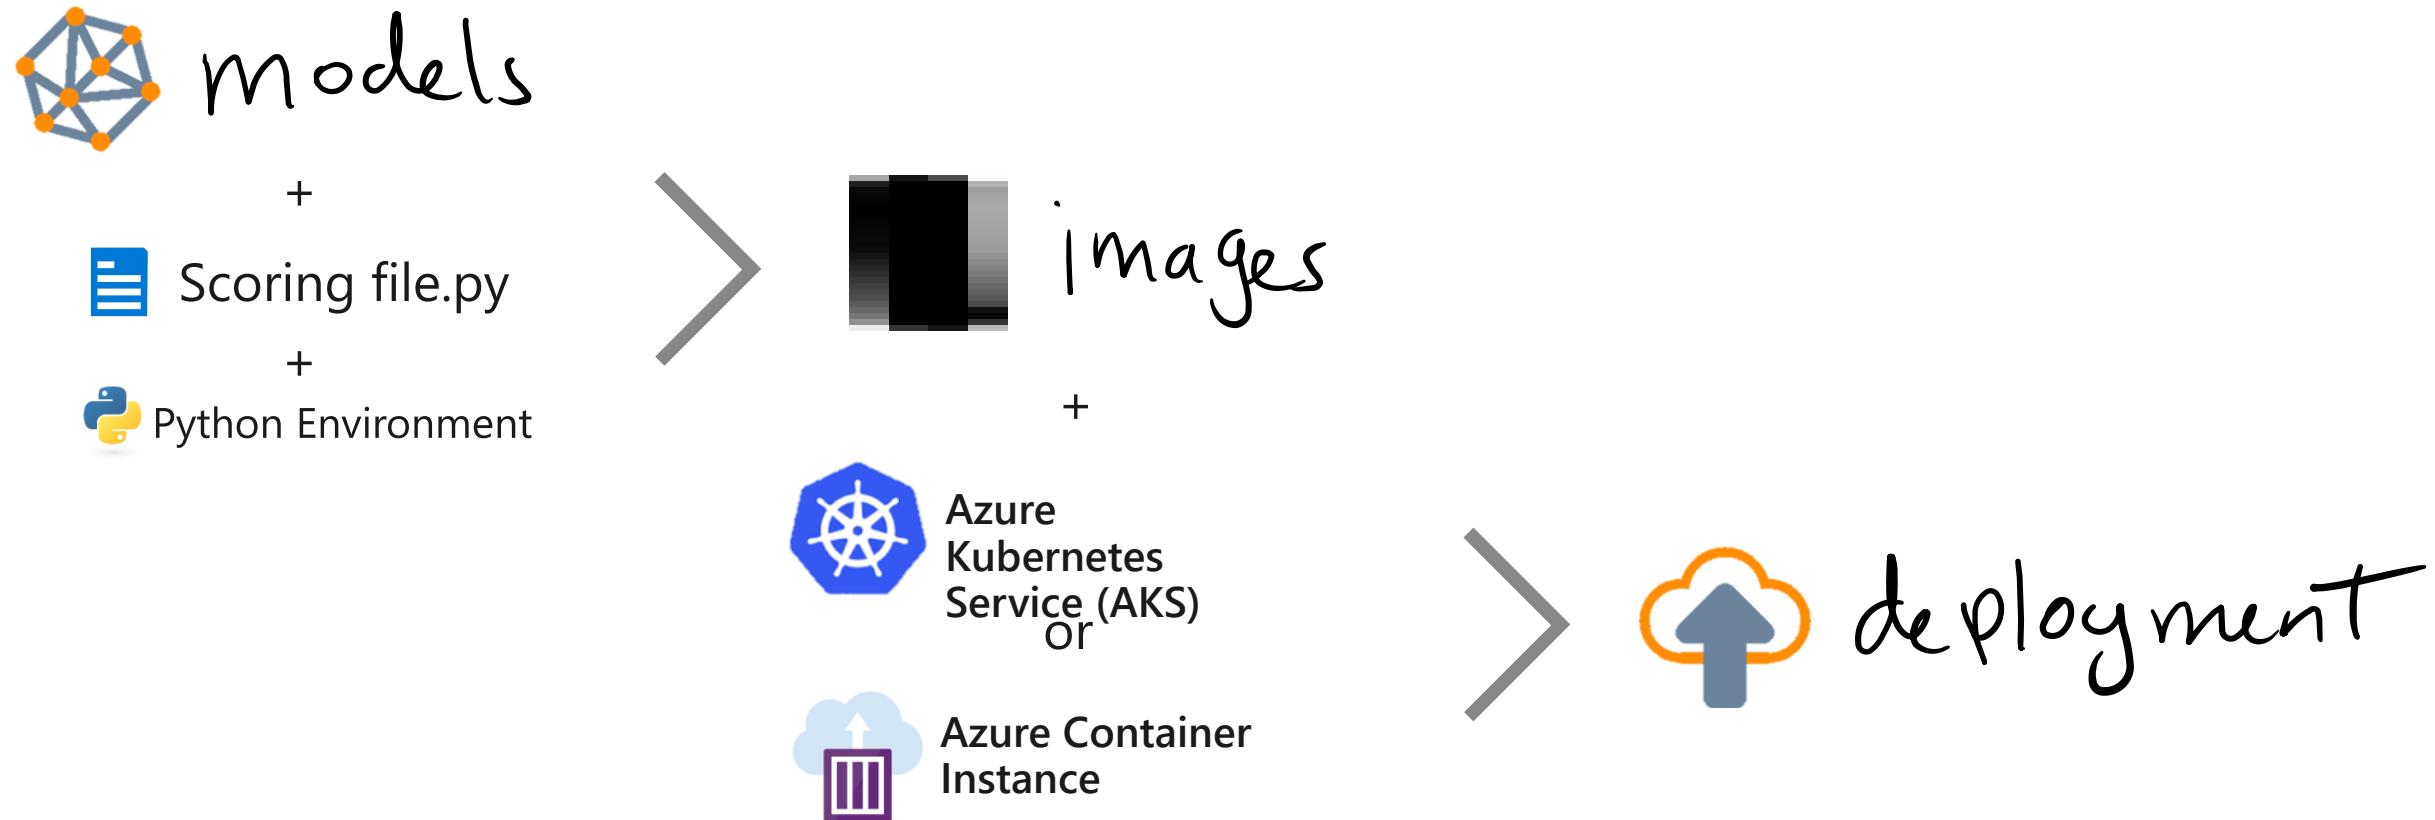
## aci-service-mnist

[Back to Deployments](#) [Edit](#) [Delete](#)

[Details](#) [Models](#) [Images](#)

Attributes	
Description	ACI deployment of MNIST Model
State	Healthy
Compute Type	ACI
Service ID	aci-service-mnist
Tags	mnist
Creation date	09/16/2018, 5:52:18 PM UTC
Last updated	09/16/2018, 5:52:27 PM UTC
Image ID	tf-mnist:1
Scoring URI	<a href="http://40.121.221.168:80(score)">http://40.121.221.168:80(score)</a> 
CPU	0.1
Memory	0.5 GB

# Deploy image

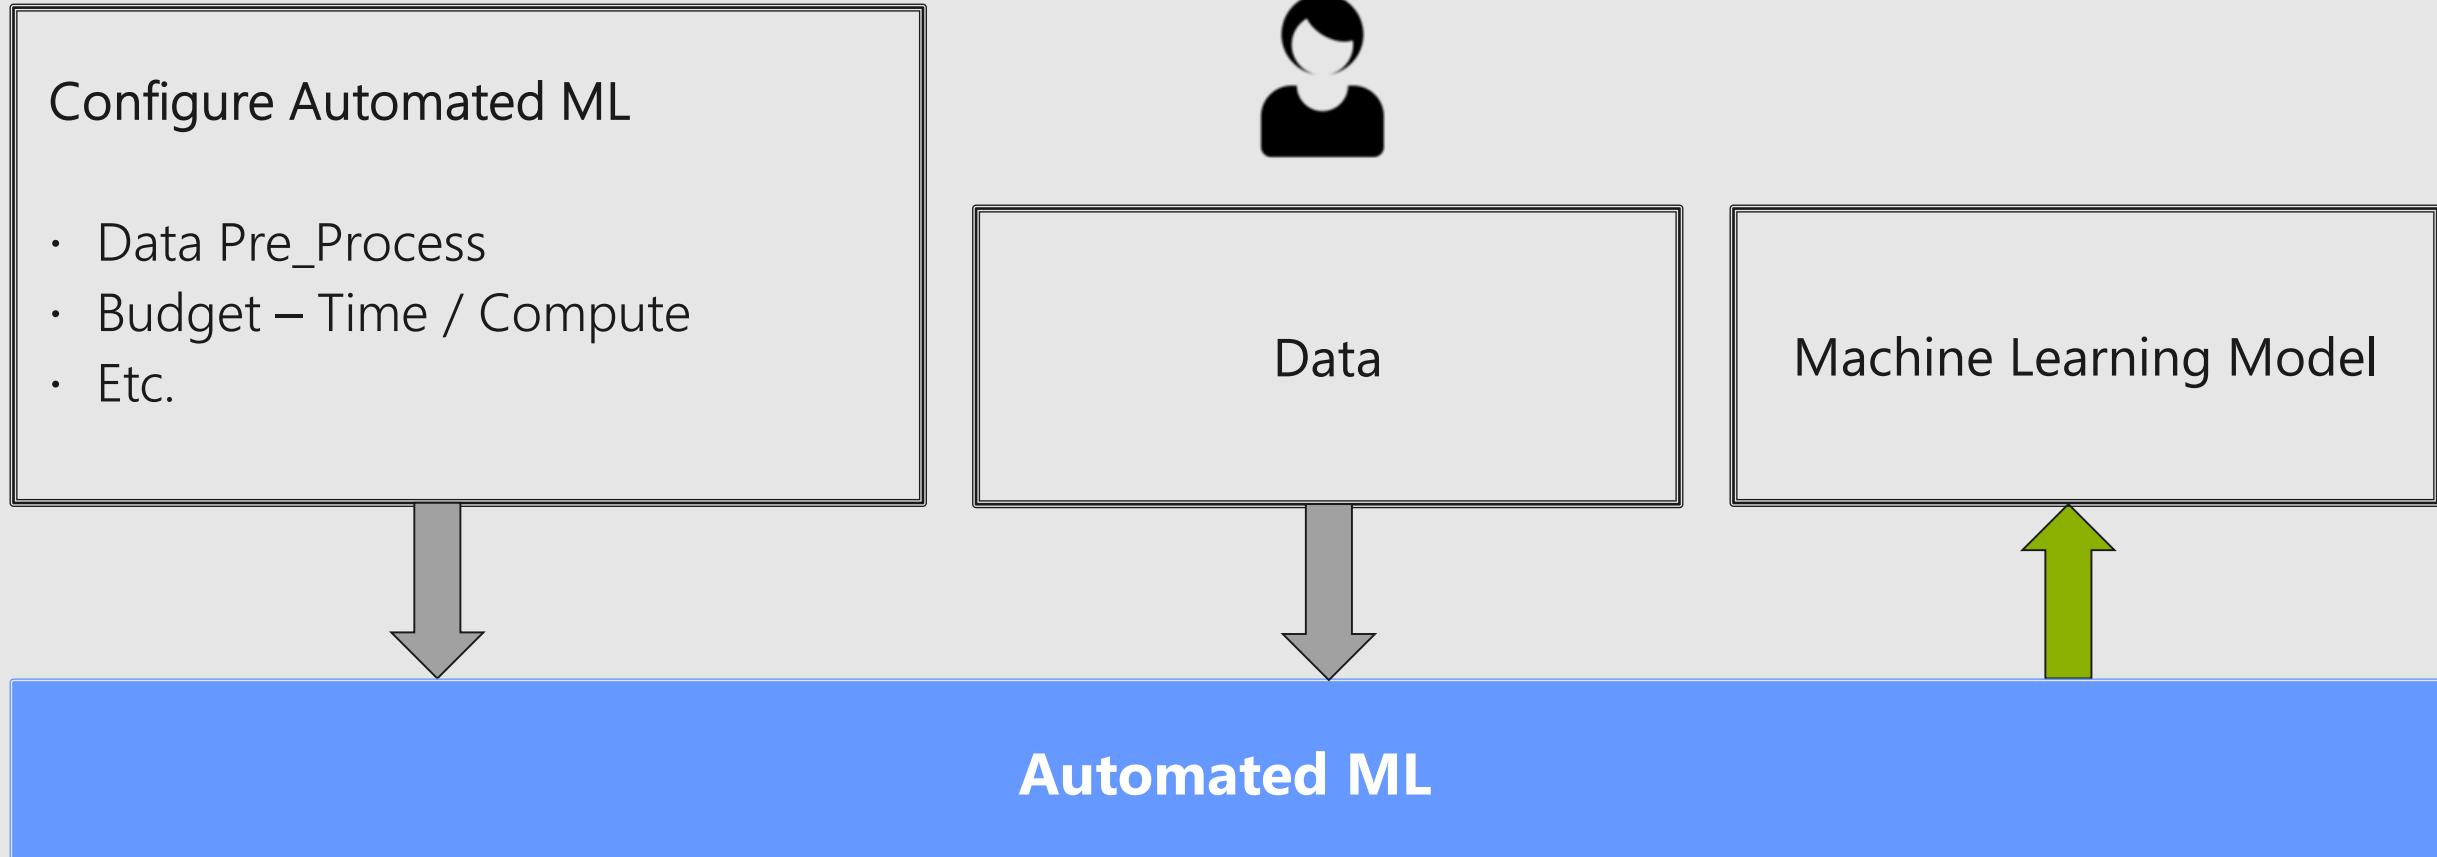


# Demo

Model Deployment

# Hyperparameter tuning

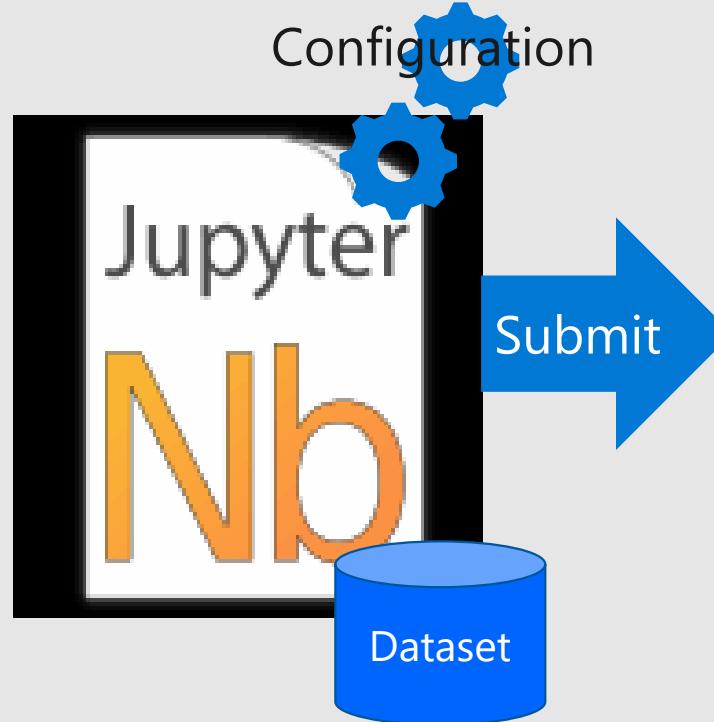
# User inputs & system outputs



# Automated ML – User Experience

Python Script

Configuration



Automated ML



Generate  
Algorithms &  
Hyperparameter  
values

Train

Models (User Compute – Local or Cloud)

Iteration	Pipeline	Iteration metric	Best metric
0	SparseNormalizer, LogisticRegression	0.99755788	0.99755788
1	StandardScalerWrapper, KNeighborsClassifier	0.99788041	0.99788041
2	MaxAbsScaler, LightGBMClassifier	0.99827043	0.99827043
3	MaxAbsScaler, DecisionTreeClassifier	0.8321242	0.99827043
4	SparseNormalizer, LightGBMClassifier	0.99794397	0.99827043
5	StandardScalerWrapper, KNeighborsClassifier	0.9983558	0.9983558
6	StandardScalerWrapper, LightGBMClassifier	0.99883517	0.99883517
7	StandardScalerWrapper, SGDClassifierWrapper	0.99455258	0.99883517
8	MaxAbsScaler, LightGBMClassifier	0.99753115	0.99883517
9	StandardScalerWrapper, KNeighborsClassifier	0.99863863	0.99883517

Pages: 1 2 3 4 Next Last 10 per page

Output

High Quality Machine Learning Model

# Easy to Use

jupyter 102.auto-ml-regression (unsaved changes) Logout

File Edit View Insert Cell Kernel Widgets Help Not Trusted Python [default]

Code

Instantiate Auto ML Regressor

Instantiate a AutoML Object This creates an Experiment in Azure ML. You can reuse this objects to trigger multiple runs. Each run will be part of the same experiment.

Property	Description
primary_metric	This is the metric that you want to optimize. Auto ML Regressor supports the following primary metrics <i>spearman_correlation</i> <i>normalized_root_mean_squared_error</i> <i>r2_score</i>
max_time_sec	Time limit in seconds for each iterations
iterations	Number of iterations. In each iteration Auto ML Classifier trains the data with a specific pipeline
num_cross_folds	Cross Validation split

In [5]:

```
from azureml.train.automl import AutoMLConfig

automl_config = AutoMLConfig(task = 'regression',
                             debug_log = 'automl_errors.log',
                             primary_metric = 'spearman_correlation',
                             max_time_sec = 12000,
                             iterations = 10,
                             n_cross_validations = 3,
                             verbosity = logging.INFO,
                             X = X,
                             y = y,
                             path=project_folder)
```

Training the Model

You can call the fit method on the AutoML instance and pass the run configuration. For Local runs the execution is synchronous. Depending on the data and number of iterations this can run for while. You will see the currently running iterations printing to the console.

fit method on Auto ML Regressor triggers the training of the model. It can be called with the following parameters

Parameter	Description
X	(sparse) array-like, shape = [n_samples, n_features]
y	(sparse) array-like, shape = [n_samples, ], [n_samples, n_classes] Multi-class targets. An indicator matrix turns on multilabel classification.
compute_target	Indicates the compute used for training. <i>local</i> indicates train on the same compute which hosts the jupyter notebook. For DSVM and Batch AI please refer to the relevant notebooks.
show_output	True/False to turn on/off console output

In [6]:

```
local_run = experiment.submit(automl_config, show_output=True)
```

Parent Run ID: AutoML\_e7ad4236e-8935-4e93-888d-1ea8310a6b22

\*\*\*\*\*  
ITERATION: The iteration being evaluated.  
PIPELINE: A summary description of the pipeline being evaluated.  
DURATION: Time taken for the current iteration.  
METRIC: The result of computing score on the fitted pipeline.  
BEST: The best observed score thus far.  
\*\*\*\*\*

ITERATION	PIPELINE	DURATION	METRIC	BEST
0	Normalize extra trees regressor	0:00:12.069893	0.688	0.688
1	Normalize lightGBM regressor	0:00:11.192919	0.597	0.688
2	Normalize Elastic net	0:00:09.866233	0.689	0.689
3	Scale 0/1 lightGBM regressor	0:00:10.069764	0.656	0.689
4	Robust Scaler KNN regressor	0:00:09.090668	0.598	0.689
5	Normalize lightGBM regressor	0:00:12.562876	0.649	0.689
6	Robust Scaler KNN regressor	0:00:09.361137	0.600	0.689
7	Normalize SGD regressor	0:00:09.910672	0.070	0.689
8	Scale 0/1 extra trees regressor	0:00:10.442752	0.685	0.689
9	Robust Scaler Gradient boosting regressor	0:00:09.567582	0.651	0.689

# Demo

AutoML on remote VM

# Demo

Hyperparameter tuning

